

Computational **Audio** and **Image** Analysis

with the **Spectrograms** library

Abstract

This manual documents the mathematical foundations and computational implementation of the **spectrograms** library, a precision-generic Rust library for FFT-based audio and image analysis. It covers the Discrete Fourier Transform and fast FFT algorithms; windowed spectral analysis via the Short-Time Fourier Transform; perceptual frequency scales including the Mel scale, Equivalent Rectangular Bandwidth, and the Constant-Q Transform; advanced audio features such as Mel-frequency cepstral coefficients, chromagrams, and binaural spatial analysis via ITD, IPD, ILD, and ILR; two-dimensional Fourier methods for image filtering and edge detection; streaming convolution via the overlap-save method; and numerical precision considerations for single and double-precision computation. All algorithms are presented in closed-form mathematical notation alongside their Rust and Python implementations. The manual assumes familiarity with linear algebra, complex numbers, and basic discrete-time signal processing.

Contents

1	Introduction	1
1.1	Target Audience and Prerequisites	1
1.2	Scope and Organization	1
1.3	Numeric Precision: Generic over f32 and f64	2
1.4	Notation Conventions	4
2	Fourier Transform Foundations	4
2.1	The Discrete Fourier Transform (DFT)	4
2.2	Fast Fourier Transform (FFT): The Cooley-Tukey Algorithm	7
2.3	Real-to-Complex FFT and Hermitian Symmetry	8
2.4	Inverse FFT and Perfect Reconstruction	10
2.5	2D FFT via Row-Column Decomposition	10
3	Window Functions for Spectral Analysis	11
3.1	The Windowing Problem	11
3.2	Windowing Convention: Symmetric vs. Periodic	12
3.3	Rectangular Window (No Windowing)	13
3.4	Hanning Window	14
3.5	Hamming Window	14
3.6	Blackman Window	15
3.7	Kaiser Window and Bessel Functions	16
3.8	Gaussian Window	17
3.9	Trade-offs: Frequency Resolution vs. Spectral Leakage	18
3.10	Constant Overlap-Add (COLA) Constraint	21
4	Short-Time Fourier Transform (STFT)	24
4.1	Time-Frequency Analysis Motivation	24
4.2	STFT Definition and Computation	25
4.3	Framing and Overlap	25

4.4	STFT Computation Pipeline	27
4.5	Spectrogram as STFT Magnitude	29
4.6	Time-Frequency Uncertainty Principle	30
5	Amplitude Scaling and Representations	32
5.1	Power Spectrum	32
5.2	Magnitude Spectrum	32
5.3	Decibel Scale	33
5.4	Parseval's Theorem and Energy Conservation	34
6	Perceptual Frequency Scales	35
6.1	Motivation: Human Auditory Perception	36
6.2	Mel Scale	36
6.3	ERB Scale: Glasberg and Moore Model	40
6.4	Logarithmic Frequency (Musical Scale)	42
6.5	Filterbank Design and Sparse Matrices	44
7	Constant-Q Transform (CQT)	45
7.1	Motivation and Q Factor	45
7.2	Kernel-Based Implementation	46
7.3	Variable-Length Windows	47
7.4	Sparsity Optimization	48
8	Advanced Audio Features	49
8.1	Mel-Frequency Cepstral Coefficients (MFCC)	50
8.2	Chroma Features for Music Analysis	52
8.3	ERB Spectrograms and Gammatone Filters	54
9	Image Processing via FFT	55
9.1	2D Convolution and the Convolution Theorem	55
9.2	Kernel Padding and Phase Shifting	56
9.3	Spatial Filtering in the Frequency Domain	57

9.4	Gaussian Blur via FFT	57
9.5	Edge Detection and High-Pass Filtering	59
10	Binaural Spectrograms	60
10.1	Spatial Hearing and Binaural Cues	60
10.2	Interaural Time Difference (ITD)	60
10.3	Interaural Phase Difference (IPD)	61
10.4	Interaural Level Difference (ILD)	61
10.5	Interaural Level Ratio (ILR)	61
10.6	Histogram Analysis and Comparison	62
10.7	Parallelisation	63
11	Modified Discrete Cosine Transform (MDCT)	64
11.1	Motivation and Codec Context	64
11.2	MDCT Definition	64
11.3	Fast Computation via the Packing Trick	64
11.4	Perfect Reconstruction and the TDAC Condition	65
11.5	Parameters	66
12	FFT Convolution and Minimum Phase	67
12.1	FFT-Based Linear Convolution	68
12.2	Spectral-Division Deconvolution	68
12.3	Streaming Overlap-Save Convolution	69
12.4	Minimum-Phase Conversion	70
13	Computational Optimizations	72
13.1	Plan-Based Computation	72
13.2	Sparse Matrix Operations	73
13.3	Design Patterns	74
13.4	Memory Layout and Cache Efficiency	75
13.5	FFTW Integration	75

14 Numerical Considerations	76
14.1 Floating-Point Precision	76
14.2 Normalization Conventions	76
14.3 Dynamic Range and Epsilon Thresholds	77
15 Conclusion and Further Reading	78
15.1 Summary of Key Concepts	78
15.2 Applications and Extensions	79
16 Python API Reference	80
16.1 Introduction	80

List of Figures

- 1 DFT analysis of a decaying three-note chord. **Top-left:** magnitude spectrum $|X[k]|$ with the four dominant frequency bins annotated. **Top-right:** time-domain waveform showing the attack transient and exponential decay. **Bottom:** the real (cosine) and imaginary (sine) components of two DFT basis functions at the peak frequencies; the inner product of $x[n]$ with each basis vector gives the corresponding complex amplitude $X[k]$ 6

- 2 Hermitian symmetry of the full DFT for a real-valued chord signal ($N = 256$, $f_s = 8$ kHz). **Left:** magnitude $|X[k]|$ is symmetric about DC—the negative-frequency half is a mirror of the positive-frequency half. **Right:** phase $\angle X[k]$ is antisymmetric. Only the $N/2 + 1$ unique bins (positive frequencies, blue) need to be stored; the red dashed curve is entirely redundant. 9

- 3 Window functions in the time domain ($N = 512$ samples). All windows taper smoothly to reduce discontinuities at the frame boundaries; the rectangular window (uniform weight) is the limiting case with no taper. 19

- 4 Frequency-domain magnitude responses of the five window functions (normalised, zero-padded to 8192 points for a smooth curve). The main lobe width and sidelobe levels reveal the fundamental leakage–resolution trade-off: wider main lobe $\Rightarrow$ lower sidelobes. The dashed horizontal line marks the -13 dB level of the rectangular window’s first sidelobe. 20

- 5 Spectral leakage on a real two-note signal with non-integer period ($f_1 = 441.7$ Hz, $f_2 = 882.3$ Hz). **Left:** the rectangular window spreads energy across many bins (leakage sidelobes clearly visible). **Right:** the Hanning window suppresses sidelobes by over 30 dB, isolating the two peaks cleanly. 21

- 6 Constant Overlap-Add (COLA) visualisation for six overlapping frames ($N = 64$, 50% hop). **Left:** Hanning window—the coloured envelopes sum (black curve) to a constant, guaranteeing perfect reconstruction. **Right:** rectangular window—also constant, confirming that any non-zero hop satisfies COLA for the rectangular case. 24

- 7 Three amplitude representations of the same spectrogram (Hanning window, $N=2048$, 50% hop, $f_s=22,050$ Hz) computed from a four-chord piano progression. **Left:** power $|X[m, k]|^2$ —extremely wide dynamic range makes quiet components invisible. **Centre:** magnitude $|X[m, k]|$ —square-root compression reveals more structure. **Right:** decibel scale $10 \log_{10} |X|^2$ —perceptually linear, with the harmonic series of each chord clearly legible. The dB representation is standard for analysis and display. 30
- 8 The time-frequency uncertainty trade-off demonstrated on the same piano chord progression at two FFT sizes. **Left** ($N=256$, **31 ms**): individual note onsets are sharply localised in time but the harmonic partials blur together vertically. **Right** ($N=4096$, **185 ms**): individual harmonics are clearly resolved in frequency but chord transitions smear horizontally. Neither window can achieve high resolution in both dimensions simultaneously—this is not an implementation limitation but a fundamental property of Fourier analysis. 31
- 9 End-to-end feature extraction pipeline applied to synthesised vowel speech (/a/ /e/ /i/ /o/ at $f_s=22\,050$ Hz). **(a)** Raw waveform showing the formant-driven amplitude modulation of each vowel. **(b)** Linear spectrogram (dB, $N=512$): the first two formants (F_1 , F_2) shift noticeably between vowels, encoding their identity. **(c)** Mel spectrogram ($M=40$ bands): perceptually weighted, emphasising the formant region below 3 kHz. **(d)** MFCCs ($C=13$): the DCT decorrelates the mel bands and compacts each vowel’s identity into the first few coefficients—the representation used directly by most speech recognition systems. 35
- 10 **Top:** mel triangular filterbank ($M=40$ filters, $f_{\min}=0$ Hz, $f_{\max}=8$ kHz). Each overlapping triangle is wider at higher frequencies (more Hz per mel), capturing the logarithmic perception of pitch. Colour indicates filter index from low (violet) to high (yellow). **Bottom:** the Hz-to-mel mapping $\text{mel}(f) = 2595 \log_{10}(1 + f/700)$, linear below 1 kHz and logarithmic above. 39
- 11 Mel-scale spectrogram (dB) of a four-chord piano progression (C maj $\rightarrow$ A min $\rightarrow$ F maj $\rightarrow$ G maj), $M=128$ mel bands. The perceptual frequency axis allocates many bands to the musically rich low-frequency region and fewer to the upper overtones, matching the logarithmic sensitivity of human pitch perception. The harmonic series of each chord is clearly visible as parallel horizontal streaks. 40

12	Gammatone auditory filters at five centre frequencies (125 Hz, 500 Hz, 1 kHz, 2 kHz, 4 kHz). Left: time-domain impulse responses (order $n=4$, bandwidth $b=1.019 \cdot \text{ERB}(f_c)$). The envelope $t^{n-1}e^{-2\pi b t}$ shortens at higher frequencies because the ERB bandwidth grows with f_c . Right: frequency responses showing the characteristic asymmetric skirt of the auditory filter—sharper on the high-frequency side, broader toward lower frequencies.	42
13	Constant-Q Transform spectrogram of a chromatic ascending arpeggio from C2 to C8 (12 bins/octave, 6 octaves, $f_{\min}=65$ Hz, dB scale). Unlike the STFT, successive semitones are uniformly spaced vertically—a direct consequence of the constant- Q design. Each note produces a brief burst that aligns precisely with its pitch-class label on the y -axis, demonstrating the transform’s suitability for pitch-accurate music analysis.	49
14	MFCCs ($C=13$) computed from five synthesised vowels (/a/, /e/, /i/, /o/, /u/, each ≈ 0.55 s). Vertical white dashes mark vowel transitions. The low-order coefficients (rows 0–4) encode the broad spectral envelope (vowel identity), while higher-order coefficients capture finer spectral detail. The DCT-II decorrelates the mel-band energies, concentrating the discriminative information into the first few coefficients—which is why 13 MFCCs are sufficient for most speech tasks.	52
15	Chromagram of the four-chord progression (C maj $\rightarrow$ A min $\rightarrow$ F maj $\rightarrow$ G maj). The y -axis lists the 12 pitch classes; energy is summed across all octaves, so the display is octave-invariant. Within the C major segment (0–1 s), the C, E, and G rows light up simultaneously; A minor activates A, C, and E; and so on. This compact 12-dimensional vector is sufficient for key detection, chord recognition, and cover-song identification.	53
16	ERB/gammatone spectrogram of the chord progression (log-power scale, 64 channels from 80 Hz to 8 kHz). Each channel’s bandwidth grows with frequency following the ERB function, providing finer frequency resolution in the low-frequency region where harmonic relationships are most salient.	55

- 17 Image processing pipeline via 2D FFT applied to a 320×320 greyscale crop of the raccoon portrait (`scipy.datasets.face`). (1) **Original**: the unmodified input. (2) **2D FFT magnitude** ($\log |F|$): the DC component (centre) is the brightest point; horizontal and vertical lines correspond to strong directional content in the image. (3) **Gaussian blur** ($\sigma=4$): low-pass filtering via multiplication by a Gaussian mask in the frequency domain. (4) **High-pass edges** ($r<20$): zeroing the central disc retains only high-frequency content, revealing structural edges. 59
- 18 Binaural spectrograms for a simulated two-source scene: a piano chord at 30° right ($\text{ITD} \approx 340 \mu\text{s}$, $\text{ILD} \approx 8 \text{ dB}$) mixed with a chirp at 60° left. All four cues are shown as time-frequency maps (red = rightward/positive, blue = leftward/negative). **ITD** is most reliable at low frequencies ($< 1.5 \text{ kHz}$); **IPD** wraps at high frequencies due to sub-wavelength ambiguity; **ILD** dominates above 1.5 kHz where the head creates an acoustic shadow; **ILR** normalises the power asymmetry to $[-1, +1]$ and is interpretable at all frequencies. 63
- 19 MDCT analysis and perfect reconstruction of a 2-second piano chord progression (sine window, $2N=2048$, 50% hop). **Top**: MDCT coefficient matrix $\log_{10} |C[k, m]|$, analogous to a spectrogram but with 50% dimensionality reduction: N input samples map to $N/2$ real coefficients per frame. **Centre**: original waveform (blue) overlaid with the IMDCT reconstruction (red dashed)—the two are visually indistinguishable. **Bottom**: reconstruction error, which is at the level of floating-point rounding noise ($\approx 10^{-14}$), confirming the Time-Domain Aliasing Cancellation (TDAC) property of the sine window. 67
- 20 Linear-phase vs. minimum-phase conversion of a 64-tap low-pass FIR filter ($f_c=2 \text{ kHz}$, $f_s=22\,050 \text{ Hz}$, Hamming window) via the real-cepstrum method. **Top-left**: impulse responses—the linear-phase filter (blue) is symmetric around its centre tap; the minimum-phase equivalent (orange) concentrates energy at the start. **Top-right**: magnitude responses—both filters share *identical* frequency responses; only the phase differs. **Bottom-left**: phase responses. **Bottom-right**: group delay—the linear-phase filter introduces $(L-1)/2 \approx 1.45 \text{ ms}$ of constant latency; the minimum-phase design reduces this to near zero, enabling low-latency room correction and live processing. 71

List of Code Listings

1	The same call at the default (f64) and at f32	3
2	R2C output size for real FFT input	8
3	Verifying C2R inverse FFT round-trip	10
4	2D FFT function signature and output shape	11
5	Generic make-window constructor signature	13
6	Rectangular window implementation	13
7	Hanning window implementation	14
8	Hamming window implementation	15
9	Blackman window implementation	15
10	Kaiser window with modified Bessel function	16
11	Gaussian window implementation	17
12	Verifying COLA constraint for window and hop	22
13	Frame count with optional centre-padding	26
14	Workspace struct with reusable STFT buffers	28
15	Windowed frame extraction and FFT per frame	28
16	Building the sparse mel filterbank matrix	37
17	Log-frequency interpolation via sparse matrix	43
18	Row-wise sparse matrix storage definition	44
19	Sparse matrix-vector multiplication	44
20	CQT kernel with windowed complex exponential	46
21	Zeroing small CQT kernel coefficients	48
22	DCT-II implementation for MFCC cepstral step	50
23	Sinusoidal liftering applied to MFCC array	51
24	Padding kernel with centre at origin for FFT	56
25	2D Gaussian kernel generation and normalisation	57
26	Computing ITD spectrogram from stereo channels	60
27	Mean ITD difference between reference and test	62
28	Enabling rayon parallelisation in Cargo.toml	63
29	MDCT and IMDCT with sine-window reconstruction	65
30	MDCT at single-precision f32 via turbofish	66
31	FFT linear convolution with NonEmptySlice input	68
32	Regularised spectral deconvolution	68
33	Streaming overlap-save with fixed block size	69
34	Converting linear-phase FIR to minimum phase	70
35	Reusable STFT plan with cached FFT and window	72
36	Sparse matrix layout and filterbank multiply	73
37	Per-frame allocation — wasteful hot-loop pattern	74
38	Pre-allocated workspace reused across frames	74
39	Full workspace struct with resizable buffers	74
40	Asserting row-major layout for FFT compatibility	75
41	FFTW planning modes from ESTIMATE to PATIENT	75
42	Log floor with epsilon to avoid negative infinity	77

43	Epsilon threshold for sparse matrix storage	77
44	Installing the spectrograms package	80
45	Importing numpy and spectrograms	80
46	Catching spectrogram-specific exceptions	81
47	Computing a 1D FFT with Hermitian output	81
48	Inverse FFT and reconstruction error check	81
49	Power and magnitude spectra with windowing	82
50	Selecting predefined window function types	83
51	Parametric Kaiser and Gaussian windows	83
52	Configuring STFT parameters and presets	84
53	Power, magnitude, and dB linear spectrograms	84
54	Accessing metadata from a Spectrogram object	85
55	Computing mel-scale spectrograms	86
56	Computing ERB-scale spectrograms	86
57	Computing log-Hz spectrograms for music	87
58	Computing CQT spectrograms	87
59	Raw complex STFT with magnitude and phase extraction	88
60	Inverse STFT reconstruction and error verification	88
61	Computing MFCCs with liftering and DCT	89
62	Computing pitch-class chromagram	89
63	Reusing a mel-dB plan for batch processing	90
64	All spectrogram plan types from SpectrogramPlanner	90
65	Frame-by-frame spectrogram computation	91
66	Streaming real-time frame extraction from buffer	91
67	Forward 2D FFT with half-spectrum output	92
68	Inverse 2D FFT and reconstruction error	92
69	2D power and magnitude spectra with log scale	93
70	FFT shift to centre the DC component	93
71	Batch 2D FFT via Fft2dPlanner plans	93
72	Gaussian blur via FFT convolution	94
73	Low-pass, high-pass, and band-pass image filters	94
74	FFT-based edge detection and image sharpening	94
75	WindowType constructors and variants	95
76	StftParams constructor signature	95
77	LogParams constructor for dB floor	95
78	SpectrogramParams constructors and presets	95
79	MelParams constructor with frequency range	95
80	ErbParams constructor with gammatone option	96
81	LogHzParams constructor with octave resolution	96
82	CqtParams constructor with sparsity threshold	96
83	ChromaParams constructor and music preset	96
84	MfccParams constructor and speech preset	97
85	All spectrogram compute function signatures	97
86	Audio feature compute function signatures	98

87	1D FFT and spectrum compute function signatures	98
88	2D FFT function signatures and return types	99
89	Image processing function signatures	99

1 Introduction

This manual provides a comprehensive mathematical and computational overview of the algorithms, optimizations, and transformations implemented in the `spectrograms` library. The focus is on understanding the theoretical foundations and practical implementations of FFT-based signal processing for audio and image analysis.

The material emphasizes the *why* and *how* of each algorithm: *why* certain design choices matter, and *how* they are implemented efficiently. The material covers both the mathematics and the practical computational considerations that make these methods work in production systems.

1.1 Target Audience and Prerequisites

This material assumes familiarity with:

- Linear algebra (matrix multiplication, vector operations, eigenvalues)
- Complex numbers and Euler’s formula
- Discrete signals and sampling theory
- Basic programming concepts and algorithm analysis
- Calculus (integrals, derivatives, series)

No prior knowledge of Fourier analysis, psychoacoustics, or advanced signal processing is assumed.

1.2 Scope and Organization

The manual is organized into major thematic sections:

1. **Fourier Transform Theory** (Section 2) — DFT, FFT algorithms, Hermitian symmetry, 2D FFT
2. **Window Functions and Spectral Analysis** (Section 3) — Hanning, Hamming, Kaiser, Blackman, Gaussian
3. **Time-Frequency Representations** (Section 4) — STFT and spectrograms
4. **Amplitude Scaling and Representations** (Section 5) — Power, magnitude, decibels, Parseval’s theorem

5. **Perceptual Frequency Scales** (Section 6) — Mel, ERB, logarithmic scales
6. **Constant-Q Transform** (Section 7) — Variable-length kernels, sparsity optimization
7. **Advanced Audio Features** (Section 8) — MFCC, chroma, gammatone filters
8. **Image Processing via FFT** (Section 9) — 2D convolution, filtering, edge detection
9. **Binaural Spectrograms** (Section 10) — ITD, IPD, ILD, ILR spatial audio features
10. **Modified Discrete Cosine Transform** (Section 11) — MDCT/IMDCT, perfect reconstruction, codec applications
11. **FFT Convolution and Minimum Phase** (Section 12) — FFT-based convolution, streaming overlap-save, minimum-phase conversion
12. **Computational Optimizations** (Section 13) — Sparse matrices, FFT plans, caching
13. **Numerical Considerations** (Section 14) — Precision, normalization, dynamic range

1.3 Numeric Precision: Generic over f32 and f64

The library is *precision-generic*: each core computation is written once and instantiated at either single (**f32**) or double (**f64**) precision. This genericity is mediated by a single sealed trait, **Sample**, which serves as the dispatch seam between the precision-agnostic algorithms and the concrete floating-point type and its FFT backend.

The Sample Trait as a Dispatch Seam

Sample is a *sealed* trait, implemented only for **f32** and **f64**. It bundles the numeric capabilities the algorithms require (**Float**, **FloatConst**, **NumAssign**, **Copy**, **Send**, **Sync**, and **'static**) together with the scalar conversions **from_f64** and **from_usize** (used to lift literal constants and index or length values into the chosen precision) and the FFT-plan constructors (**plan_r2c**, **plan_c2r**, **plan_c2c**, and their two-dimensional counterparts). The **realf** and **FFTW** backend plan traits (**R2cPlan**, **C2cPlan**, **C2rPlan**, and their 2D variants) are themselves generic over the scalar type, so a single generic code path drives both precisions without duplication.

Default to f64 for backward compatibility. The scalar type parameter *defaults* to `f64`. Consequently every `f64` signature and example in this manual remains valid and unchanged in meaning: double precision is selected automatically when the inputs are `f64`, or when no type is specified. Single precision is strictly opt-in—it is inferred from `f32` inputs, or requested explicitly with *turbofish* syntax:

Listing 1: The same call at the default (`f64`) and at `f32`

```
// Double precision (default): T is inferred as f64
let s = stft(&sig_f64, n_fft, hop, WindowType::Hanning, true)
    );

// Single precision: inferred from f32 inputs ...
let s = stft(&sig_f32, n_fft, hop, WindowType::Hanning, true)
    );

// ... or requested explicitly via turbofish
let s = stft::<f32>(&sig_f32, n_fft, hop, WindowType::
    Hanning, true)?;
```

Precision genericity spans the full feature set: the STFT, the mel, ERB, and CQT perceptual transforms, MFCCs, chroma, the MDCT, FFT convolution, minimum-phase reconstruction, the binaural features, and the 2D FFT and image operations all operate at both precisions. The legacy `mdct_f32` and `imdct_f32` entry points have been *removed*; use the generic `mdct::<f32>` and `imdct::<f32>` instead.

Caveat: f64-Returning Convenience Constructors

A few high-level convenience constructors that accept `f64` samples—`chromagram`, `mfcc`, and `gaussian_kernel_2d`—deliberately return `f64` results. They are ergonomic entry points for the common double-precision case and are *not* generic. Single precision for these features is reached instead through the input-generic variants (`chromagram_from_spectrogram` and `mfcc_from_log_mel`) or through the low-level primitives, all of which are generic over the scalar type.

Motivation. Supporting `f32` alongside `f64` serves three concrete goals. First, many spectrogram workloads are *memory-bound*: halving the width of every sample and spectral coefficient improves cache utilisation and memory bandwidth. Second, machine-learning pipelines train in single precision, so producing features directly in `f32` avoids a conversion at the pipeline boundary. Third, modern hardware—both SIMD-capable CPUs and GPUs—processes single-precision data at higher throughput, so computing features in `f32` reduces

wall-clock time for large-scale batch workloads. Double precision nonetheless remains the default, as it is the right choice for offline analysis where numerical headroom matters (Section 14.1).

1.4 Notation Conventions

Throughout this manual:

- N denotes signal length or FFT size
- k denotes frequency bin index ($0 \leq k < N$)
- n denotes time/sample index ($0 \leq n < N$)
- $j = \sqrt{-1}$ (the imaginary unit)
- $X[k]$ denotes frequency domain values
- $x[n]$ denotes time domain values
- Vectors are bold: $\mathbf{x}$
- Complex conjugate: z^*

2 Fourier Transform Foundations

2.1 The Discrete Fourier Transform (DFT)

The ***Discrete Fourier Transform*** (DFT) decomposes a finite-length sequence into its constituent frequencies. For a sequence $x[n]$ of length N , the DFT is defined as [19]:

$$X[k] = \sum_{n=0}^{N-1} x[n] \cdot e^{-j2\pi kn/N}, \quad k = 0, 1, \dots, N-1 \quad (1)$$

where $X[k]$ represents the complex amplitude at frequency bin k . Using Euler's formula $e^{j\theta} = \cos \theta + j \sin \theta$, Equation (1) can be expanded as:

$$X[k] = \sum_{n=0}^{N-1} x[n] \left[\cos \left(\frac{2\pi kn}{N} \right) - j \sin \left(\frac{2\pi kn}{N} \right) \right] \quad (2)$$

This reveals that the DFT correlates the input signal with complex sinusoids at each frequency. The computational complexity of this direct calculation is $\mathcal{O}(N^2)$ — each of N outputs requires N multiplications.

The *inverse DFT* (IDFT) reconstructs the time-domain signal:

$$x[n] = \frac{1}{N} \sum_{k=0}^{N-1} X[k] \cdot e^{j2\pi kn/N}, \quad n = 0, 1, \dots, N-1 \quad (3)$$

The normalization factor $1/N$ ensures *perfect reconstruction*:

$$\text{IDFT}(\text{DFT}(x)) = x.$$

DFT Scaling Convention

Different implementations use different scaling conventions for the DFT and inverse DFT. The three most common conventions are:

1. **No forward scaling (this library):** Forward DFT uses coefficient 1, inverse DFT uses $1/N$
2. **Symmetric scaling:** Both forward and inverse use $1/\sqrt{N}$ (unitary transform)
3. **No inverse scaling:** Forward uses $1/N$, inverse uses 1

This library uses Convention 1: The forward DFT has *no normalization factor* (coefficient of 1), and the inverse DFT includes the $1/N$ factor. This convention is standard in most signal processing literature and libraries (FFTW, NumPy, SciPy).

Energy implications: With this scaling, Parseval's theorem (Section 5.4) takes the form:

$$\sum_{n=0}^{N-1} |x[n]|^2 = \frac{1}{N} \sum_{k=0}^{N-1} |X[k]|^2 \quad (4)$$

When comparing spectral amplitudes between implementations, *always verify the scaling convention* to ensure energy-consistent analysis.




Figure 1: DFT analysis of a decaying three-note chord. **Top-left:** magnitude spectrum $|X[k]|$ with the four dominant frequency bins annotated. **Top-right:** time-domain waveform showing the attack transient and exponential decay. **Bottom:** the real (cosine) and imaginary (sine) components of two DFT basis functions at the peak frequencies; the inner product of $x[n]$ with each basis vector gives the corresponding complex amplitude $X[k]$.

2.2 Fast Fourier Transform (FFT): The Cooley-Tukey Algorithm

The **Fast Fourier Transform** (FFT) is not a different transform, but an efficient algorithm for computing the DFT. The Cooley-Tukey algorithm [7], published in 1965 (though similar ideas date to Gauss in 1805 [13]), reduces complexity from $\mathcal{O}(N^2)$ to $\mathcal{O}(N \log N)$.

The key insight is divide-and-conquer using the *symmetry* and *periodicity* properties of complex exponentials:

Twiddle Factor Properties

Let $W_N^k = e^{-j2\pi k/N}$ be the **twiddle factor**. These complex exponentials satisfy:

$$W_N^{k+N/2} = -W_N^k \quad (\text{symmetry}) \quad (5)$$

$$W_N^{k+N} = W_N^k \quad (\text{periodicity}) \quad (6)$$

Periodicity means the complex exponential repeats every N samples.

Symmetry means that values half an FFT bin apart differ only in sign.

These properties enable the FFT's recursive decomposition.

For N even, split the DFT into even and odd indexed samples:

$$X[k] = \sum_{n=0}^{N/2-1} x[2n]W_N^{2nk} + \sum_{n=0}^{N/2-1} x[2n+1]W_N^{(2n+1)k} \quad (7)$$

Using $W_N^{2k} = W_{N/2}^k$, this becomes:

$$X[k] = \underbrace{\sum_{n=0}^{N/2-1} x[2n]W_{N/2}^{nk}}_{E[k]} + W_N^k \underbrace{\sum_{n=0}^{N/2-1} x[2n+1]W_{N/2}^{nk}}_{O[k]} \quad (8)$$

where $E[k]$ and $O[k]$ are $N/2$ -point DFTs of the even and odd samples. This recursive splitting continues until the base cases are reached (typically $N = 1$ or small prime factors).

The complexity analysis: At each of $\log_2 N$ levels, N operations (additions and a twiddle factor multiplication) are performed, yielding $\mathcal{O}(N \log N)$ total complexity.

2.3 Real-to-Complex FFT and Hermitian Symmetry

When the input signal $x[n]$ is real-valued (as in audio processing), the DFT output exhibits a special property:

Hermitian Symmetry

For a real-valued signal $x[n]$, the DFT satisfies **Hermitian symmetry**:

$$X[k] = X^*[N - k] \quad (9)$$

where X^* denotes complex conjugate. This means:

- $\text{Re}(X[k]) = \text{Re}(X[N - k])$ (real parts are symmetric)
- $\text{Im}(X[k]) = -\text{Im}(X[N - k])$ (imaginary parts are antisymmetric)
- $X[0]$ and $X[N/2]$ (for even N) are purely real

This property follows directly from the DFT definition and the fact that $e^{j\theta} + e^{-j\theta} = 2\cos\theta$ is real for real θ .

Critical constraint for R2C transforms: For even N , the Nyquist frequency bin $X[N/2]$ *must be purely real* (i.e., $\text{Im}(X[N/2]) = 0$). This follows from Hermitian symmetry since $X[N/2] = X^*[N - N/2] = X^*[N/2]$, which requires the imaginary part to be zero. Violating this constraint results in non-real values after inverse FFT, causing numerical errors and incorrect signal reconstruction.

This symmetry makes it sufficient to store only the positive frequencies:

$$k \in \{0, 1, 2, \dots, N/2\} \quad (10)$$

This yields $N/2 + 1$ complex values instead of N , halving storage requirements. The library computes real-to-complex (R2C) FFTs with output size:

$$\text{output_size} = \lfloor N/2 \rfloor + 1 \quad (11)$$

Implementation detail: The library uses specialized R2C FFT algorithms (from RealFFT or FFTW backends) that exploit this symmetry for both speed and memory efficiency.

Listing 2: R2C output size for real FFT input

```
pub const fn r2c_output_size(n: usize) -> usize {  
    n / 2 + 1  
}
```

```
}
```

```
// Example: 512-point FFT on real signal
let samples = vec![0.0; 512];
let spectrum = fft(&samples, 512)?;
assert_eq!(spectrum.len(), 257); // 512/2 + 1
```




Figure 2: Hermitian symmetry of the full DFT for a real-valued chord signal ($N = 256$, $f_s = 8$ kHz). **Left:** magnitude $|X[k]|$ is symmetric about DC—the negative-frequency half is a mirror of the positive-frequency half. **Right:** phase $\angle X[k]$ is antisymmetric. Only the $N/2 + 1$ unique bins (positive frequencies, blue) need to be stored; the red dashed curve is entirely redundant.

2.4 Inverse FFT and Perfect Reconstruction

The inverse FFT (IFFT) transforms frequency-domain data back to the time domain. For a complex-to-real (C2R) inverse FFT, given Hermitian-symmetric input $X[k]$:

$$x[n] = \frac{1}{N} \sum_{k=0}^{N-1} X[k] e^{j2\pi kn/N} \quad (12)$$

Due to Hermitian symmetry, it is sufficient to sum over the non-redundant frequencies:

$$x[n] = \frac{1}{N} \left[X[0] + 2 \sum_{k=1}^{N/2-1} \operatorname{Re}(X[k] e^{j2\pi kn/N}) + X[N/2] \cos(\pi n) \right] \quad (13)$$

The library's C2R inverse FFT ensures perfect reconstruction:

Listing 3: Verifying C2R inverse FFT round-trip

```
let original = vec![1.0, 2.0, 3.0, 4.0];
let spectrum = fft2d(&original_image)?;
let reconstructed = ifft2d(&spectrum, original_image.ncols())
    ?;
// reconstructed approximately original (within floating-
// point precision)
```

2.5 2D FFT via Row-Column Decomposition

For 2D data like images with dimensions $M \times N$, the 2D DFT is:

$$X[k_1, k_2] = \sum_{n_1=0}^{M-1} \sum_{n_2=0}^{N-1} x[n_1, n_2] \cdot e^{-j2\pi(k_1 n_1/M + k_2 n_2/N)} \quad (14)$$

The separability property allows this to be decomposed into row and column transforms:

$$X[k_1, k_2] = \sum_{n_1=0}^{M-1} e^{-j2\pi k_1 n_1/M} \underbrace{\left[\sum_{n_2=0}^{N-1} x[n_1, n_2] \cdot e^{-j2\pi k_2 n_2/N} \right]}_{\text{1D FFT along each row}} \quad (15)$$

Algorithm:

1. Apply 1D FFT to each row: M transforms of length N
2. Apply 1D FFT to each column of the result: N transforms of length M

For a real-valued $M \times N$ image, the output is $M \times (N/2 + 1)$ due to Hermitian symmetry along rows.

Implementation: The library uses row-major layout and ensures contiguous memory:

Listing 4: 2D FFT function signature and output shape

```
pub fn fft2d<T: Sample>(data: &ArrayView2<T>) ->
    SpectrogramResult<Array2<Complex<T>>> {
    let (nrows, ncols) = (data.nrows(), data.ncols());

    if !data.is_standard_layout() {
        return Err(SpectrogramError::invalid_input(
            "array must be contiguous and row-major"));
    }

    let out_shape = r2c_output_size_2d(nrows, ncols);
    // out_shape = (nrows, ncols/2 + 1)

    let mut plan = planner.plan_r2c_2d(nrows, ncols)?;
    // ... perform row-column FFT
}
```

Complexity: $O(MN \log N + MN \log M) = O(MN \log(MN))$ for square images.

3 Window Functions for Spectral Analysis

3.1 The Windowing Problem

Spectral Leakage

When the DFT of a finite signal segment is computed, the signal is implicitly multiplied by a rectangular window (ones inside the segment, zeros outside). This abrupt truncation causes *spectral leakage*: energy from one frequency “leaks” into adjacent frequency bins.

Consider a pure sinusoid at frequency f_0 . If sampling captures exactly an integer number of periods, the DFT produces a perfect spike at f_0 . But with non-integer periods, the energy spreads across many bins. This occurs because multiplying by a rectangular window in the time domain is equivalent to convolving with a sinc function in the frequency domain, which has slowly-decaying sidelobes.

The fundamental trade-off [15]:

- **Narrow main lobe** $\rightarrow$ better frequency resolution
- **Low sidelobes** $\rightarrow$ less spectral leakage

Window functions smooth the transition to zero at the segment boundaries, reducing leakage at the cost of frequency resolution.

3.2 Windowing Convention: Symmetric vs. Periodic

Symmetric and Periodic Windows

Window functions can be defined with two distinct conventions:

Symmetric windows use $N - 1$ in the denominator:

$$w_{\text{sym}}[n] = f\left(\frac{n}{N-1}\right), \quad n \in [0, N-1] \quad (16)$$

Periodic windows use N in the denominator:

$$w_{\text{per}}[n] = f\left(\frac{n}{N}\right), \quad n \in [0, N-1] \quad (17)$$

The difference is subtle but important:

- **Symmetric:** Values at endpoints match (e.g., both zero for Hanning). Standard for spectral analysis.
- **Periodic:** The window is periodic with period N , so $w[N] = w[0]$. Required for perfect reconstruction in overlap-add systems.

For the Constant Overlap-Add (COLA) property [15]:

$$\sum_{m=-\infty}^{\infty} w[n - mH] = C \quad (\text{constant}) \quad (18)$$

where H is the hop size. Periodic windows satisfy COLA for specific overlap ratios, enabling perfect signal reconstruction.

Library convention: All window implementations in this library use the *symmetric* formulation ($N - 1$), which is standard for analysis applications. For synthesis requiring COLA, use appropriate overlap ratios (e.g., 50% for Hanning).

3.3 Rectangular Window (No Windowing)

The rectangular window is simply:

$$w[n] = 1, \quad n \in [0, N - 1] \quad (19)$$

Its frequency response has the narrowest main lobe but highest sidelobes (≈ -13 dB), causing severe spectral leakage.

Generic over precision. The window constructors are generic over the scalar type. The match-arm listings throughout this section show the `f64` instantiation, whereas the public entry point has the generic signature

Listing 5: Generic make-window constructor signature

```
pub fn make_window<T: Sample>(window: WindowType,
                               n_fft: NonZeroUsize) ->
    NonEmptyVec<T>
```

and the numeric literals (e.g. 0.5, 0.54) are lifted to the chosen precision through `T::from_f64` in the generic code.

Listing 6: Rectangular window implementation

```
WindowType::Rectangular => {
    w.fill(1.0);
}
```

Use when: You need maximum frequency resolution and know the signal contains exact integer periods.

3.4 Hanning Window

Also called the Hann window (named after Julius von Hann, not Richard Hamming [15]), defined as:

$$w[n] = 0.5 - 0.5 \cos\left(\frac{2\pi n}{N-1}\right), \quad n \in [0, N-1] \quad (20)$$

This is a raised cosine function. It smoothly tapers to zero at both endpoints.

Properties:

- Main lobe width: 4 bins (2× rectangular)
- First sidelobe: ≈ -32 dB
- Sidelobe falloff: -18 dB/octave

Listing 7: Hanning window implementation

```
WindowType::Hanning => {  
    let n1 = (n_fft - 1) as f64;  
    for (n, v) in w.iter_mut().enumerate() {  
        *v = 0.5 - 0.5 * cos(2.0 * PI * (n as f64) / n1);  
    }  
}
```

Why this works: The cosine taper reduces discontinuities at the edges. In the frequency domain, the Hanning window's transform is the convolution of the rectangular window's transform with two delta functions, producing lower sidelobes.

The Hanning window is the most commonly used for general-purpose spectral analysis.

3.5 Hamming Window

Named after Richard Hamming [15], this window uses optimized coefficients:

$$w[n] = 0.54 - 0.46 \cos\left(\frac{2\pi n}{N-1}\right) \quad (21)$$

The coefficients (0.54, 0.46) were chosen to minimize the first sidelobe:

- Main lobe width: 4 bins (same as Hanning)

- First sidelobe: ≈ -43 dB (better than Hanning)
- Sidelobe falloff: -6 dB/octave (worse than Hanning)

Listing 8: Hamming window implementation

```
WindowType::Hamming => {
    let n1 = (n_fft - 1) as f64;
    for (n, v) in w.iter_mut().enumerate() {
        *v = 0.54 - 0.46 * cos(2.0 * PI * (n as f64) / n1);
    }
}
```

Trade-off: The Hamming window doesn't reach exactly zero at the endpoints (it's 0.08), which can cause issues for overlap-add reconstruction. However, its lower first sidelobe makes it useful when strong interfering signals exist.

3.6 Blackman Window

A three-term cosine sum [3] with excellent sidelobe suppression:

$$w[n] = 0.42 - 0.5 \cos\left(\frac{2\pi n}{N-1}\right) + 0.08 \cos\left(\frac{4\pi n}{N-1}\right) \quad (22)$$

Properties:

- Main lobe width: 6 bins ($3\times$ rectangular)
- First sidelobe: ≈ -58 dB
- Sidelobe falloff: -18 dB/octave

Listing 9: Blackman window implementation

```
WindowType::Blackman => {
    let n1 = (n_fft - 1) as f64;
    for (n, v) in w.iter_mut().enumerate() {
        let a = 2.0 * PI * (n as f64) / n1;
        *v = 0.42 - 0.5 * cos(a) + 0.08 * cos(2.0 * a);
    }
}
```

Use when: You need minimal spectral leakage and can tolerate wider main lobes (reduced frequency resolution).

3.7 Kaiser Window and Bessel Functions

The Kaiser window [16] provides a tunable parameter β that controls the trade-off between main lobe width and sidelobe level:

$$w[n] = \frac{I_0\left(\beta\sqrt{1 - \left(\frac{2n}{N-1} - 1\right)^2}\right)}{I_0(\beta)} \quad (23)$$

Modified Bessel Function

The *modified Bessel function of the first kind*, order zero, $I_0(x)$ is defined by:

$$I_0(x) = \sum_{k=0}^{\infty} \frac{1}{(k!)^2} \left(\frac{x}{2}\right)^{2k} \quad (24)$$

For computational efficiency, a truncated power series is used. The series converges rapidly; typically 20–30 terms achieve machine-precision accuracy for $x < 10$.

Parameter β effects:

- $\beta = 0$: rectangular window
- $\beta = 5$: similar to Hamming
- $\beta = 8.6$: similar to Blackman
- Larger β : lower sidelobes, wider main lobe

Listing 10: Kaiser window with modified Bessel function

```
fn besseli0(x: f64) -> f64 {
    let mut sum = 1.0;           // k=0 term
    let mut term = 1.0;
    let half_x = x / 2.0;

    // Truncated series: sufficient for Kaiser windows
    for k in 1..30 {
        term *= (half_x / k as f64).powi(2);
        sum += term;

        // Early termination if converged
        if term < 1e-12 * sum {
            break;
        }
    }
}
```

```

    }
}

sum
}

WindowType::Kaiser { beta } => {
    let denominator = bessell_i0(beta);

    for i in 0..n_fft {
        let n = i as f64;
        let n_max = (n_fft - 1) as f64;
        let alpha = (n - n_max / 2.0) / (n_max / 2.0);
        let bessell_arg = beta * (1.0 - alpha * alpha).sqrt()
    ;

        w[i] = bessell_i0(bessell_arg) / denominator;
    }
}

```

The Kaiser window is named after James Kaiser who introduced it in 1966 for optimal FIR filter design.

3.8 Gaussian Window

A Gaussian function centered at the window midpoint:

$$w[n] = \exp\left(-\frac{1}{2}\left(\frac{n - \frac{N-1}{2}}{\sigma}\right)^2\right) \quad (25)$$

where σ (standard deviation) controls the width. Smaller $\sigma \rightarrow$ narrower window $\rightarrow$ more time localization.

Listing 11: Gaussian window implementation

```

WindowType::Gaussian { std } => {
    for i in 0..n_fft {
        let n = i as f64;
        let center = (n_fft - 1) as f64 / 2.0;
        let exponent = -0.5 * ((n - center) / std).powi(2);
        w[i] = exp(exponent);
    }
}

```

The Gaussian window has a unique property: its Fourier transform is also Gaussian. This gives optimal time-frequency localization per the uncertainty principle:

$$\Delta t \cdot \Delta f \geq \frac{1}{4\pi} \quad (26)$$

Use for: Time-frequency analysis where minimizing uncertainty is critical (e.g., Gabor transforms).

3.9 Trade-offs: Frequency Resolution vs. Spectral Leakage

Comparison table for $N = 512$:

Window	Main Lobe	Peak Sidelobe	Falloff Rate
Rectangular	2 bins	-13 dB	-6 dB/oct
Hanning	4 bins	-32 dB	-18 dB/oct
Hamming	4 bins	-43 dB	-6 dB/oct
Blackman	6 bins	-58 dB	-18 dB/oct
Kaiser ($\beta = 8.6$)	6 bins	-60 dB	-6 dB/oct

Selection guidelines:

- General purpose, good balance: Hanning
- Need low first sidelobe: Hamming
- Need very low sidelobes: Blackman
- Custom trade-off: Kaiser with tuned β
- Optimal uncertainty: Gaussian with appropriate σ




Figure 3: Window functions in the time domain ($N = 512$ samples). All windows taper smoothly to reduce discontinuities at the frame boundaries; the rectangular window (uniform weight) is the limiting case with no taper.




Figure 4: Frequency-domain magnitude responses of the five window functions (normalised, zero-padded to 8192 points for a smooth curve). The main lobe width and sidelobe levels reveal the fundamental leakage–resolution trade-off: wider main lobe $\Rightarrow$ lower sidelobes. The dashed horizontal line marks the -13 dB level of the rectangular window’s first sidelobe.




Figure 5: Spectral leakage on a real two-note signal with non-integer period ($f_1 = 441.7$ Hz, $f_2 = 882.3$ Hz). **Left:** the rectangular window spreads energy across many bins (leakage sidelobes clearly visible). **Right:** the Hanning window suppresses sidelobes by over 30 dB, isolating the two peaks cleanly.

3.10 Constant Overlap-Add (COLA) Constraint

Perfect Reconstruction and COLA

When using the STFT for signal modification or synthesis (e.g., time-stretching, pitch-shifting, noise reduction), the time-domain signal must be reconstructed from modified spectral frames. The **Constant Overlap-Add** (COLA) constraint ensures perfect reconstruction is possible [15].

For a window function $w[n]$ and hop size H , the COLA property states that the sum of overlapping windows must be constant:

$$\sum_{m=-\infty}^{\infty} w[n - mH] = C \quad (\text{constant for all } n) \quad (27)$$

where C is a non-zero constant (typically normalized to 1).

Physical interpretation: When frames are overlapped by hop size H and summed (overlap-add), every sample in the reconstructed signal

receives equal weighting from the windows. This prevents amplitude modulation artifacts.

COLA conditions for common windows:

1. **Rectangular window:** COLA for any hop size $H \leq N$ (no overlap required)
2. **Hanning window (periodic):** COLA for $H = N/2, N/4, N/8, \dots$ (50%, 75%, 87.5% overlap, etc.)
 - 50% overlap ($H = N/2$): Most common, $C = 1$
 - 75% overlap ($H = N/4$): Higher redundancy, smoother synthesis
3. **Hamming window:** COLA for $H \approx 0.375N$ to $0.5N$ (requires careful tuning; 50% overlap gives $C \approx 1.08$)
4. **Blackman window:** COLA for $H \leq N/3$ (at least 67% overlap required)
5. **Kaiser window:** COLA depends on β ; generally requires 50 – 75% overlap

Verification: To check if a window/hop combination satisfies COLA, compute:

Listing 12: Verifying COLA constraint for window and hop

```
fn verify_cola(window: &[f64], hop_size: usize) -> bool {
    let n = window.len();
    let n_overlaps = (n + hop_size - 1) / hop_size;

    // Check multiple offset positions
    for offset in 0..hop_size {
        let mut sum = 0.0;

        for m in 0..n_overlaps {
            let idx = offset + m * hop_size;
            if idx < n {
                sum += window[idx];
            }
        }

        // All positions should have same sum (within
        tolerance)
        if (sum - window[0]).abs() > 1e-10 {
            return false;
        }
    }
}
```

```
}      true
```

Implications for spectrogram inversion:

When using the STFT for analysis-modification-synthesis:

1. Choose a window/hop combination that satisfies COLA
2. Apply STFT to get frames $X[m, k]$
3. Modify the spectrogram (e.g., filter, denoise, time-stretch)
4. Apply inverse FFT to each modified frame
5. Overlap-add the frames with hop size H
6. Divide by the COLA constant C (if $C \neq 1$)

The reconstructed signal $\hat{x}[n]$ will equal the original $x[n]$ if no modifications were made (perfect reconstruction).

Library convention: For analysis-only applications (spectrograms, feature extraction), COLA is not required. However, this library uses 50% overlap ($H = N/2$) by default, which satisfies COLA for Hanning windows and enables optional synthesis workflows.




Figure 6: Constant Overlap-Add (COLA) visualisation for six overlapping frames ($N = 64$, 50% hop). **Left:** Hanning window—the coloured envelopes sum (black curve) to a constant, guaranteeing perfect reconstruction. **Right:** rectangular window—also constant, confirming that any non-zero hop satisfies COLA for the rectangular case.

4 Short-Time Fourier Transform (STFT)

4.1 Time-Frequency Analysis Motivation

The standard DFT (Equation (1)) provides frequency content but discards all temporal information: it indicates *which* frequencies are present, but not *when* they occur. For non-stationary signals (speech, music, transient events), time-frequency analysis is required [2].

The **Short-Time Fourier Transform** (STFT) solves this by applying the DFT to short, overlapping segments (frames) of the signal, creating a time-frequency representation.

4.2 STFT Definition and Computation

For a signal $x[n]$ with window function $w[n]$ of length N , the STFT is:

$$X[m, k] = \sum_{n=0}^{N-1} x[n + mH] \cdot w[n] \cdot e^{-j2\pi kn/N} \quad (28)$$

where:

- m is the frame index
- H is the **hop size** (samples between consecutive frames)
- k is the frequency bin
- N is the FFT size (window length)

The result is a 2D complex matrix: $X[m, k]$ with dimensions (time frames) $\times$ (frequency bins).

4.3 Framing and Overlap

The signal is divided into overlapping frames:

$$\text{Frame } m \text{ starts at sample: } mH \quad (29)$$

Key parameters:

- **FFT size (N):** Window length, determines frequency resolution: $\Delta f = f_s/N$
- **Hop size (H):** Samples between frames, determines time resolution: $\Delta t = H/f_s$
- **Overlap:** $(N - H)/N$. Common: 50% ($H = N/2$) or 75% ($H = N/4$)

Number of frames for signal of length L samples:

$$M = \left\lfloor \frac{L - N}{H} \right\rfloor + 1 \quad (30)$$

With centering (padding $N/2$ zeros on each end):

$$M = \left\lfloor \frac{L + N - H}{H} \right\rfloor \quad (31)$$

STFT Centering and Temporal Alignment

Why pad $N/2$ zeros? Without centering, the first window is positioned with its *left edge* at sample 0. This means the window's center is at sample $N/2$, causing a temporal shift in the resulting spectrogram. By padding $N/2$ zeros *before* the signal, the first window's *center* aligns with sample 0. This ensures:

- Frame m is centered at time $t = mH/f_s$ (no offset)
- Edge effects are symmetric at signal boundaries
- Improved temporal localization for transient events

The same $N/2$ padding is applied at the signal's end, ensuring symmetric treatment of boundaries. This is standard practice in audio analysis libraries (e.g., librosa's `center=True`).

Padding type specification: This library uses *zero-padding* (also called zero-extension) for both the initial and final $N/2$ samples:

$$\tilde{x}[n] = \begin{cases} 0 & n < 0 \\ x[n] & 0 \leq n < L \\ 0 & n \geq L \end{cases} \quad (32)$$

where L is the original signal length.

Alternative: Reflective padding extends the signal by mirroring boundary values (e.g., $\tilde{x}[-1] = x[0]$, $\tilde{x}[-2] = x[1]$), which can reduce edge discontinuities for certain signals. However, zero-padding is the standard convention in most audio processing libraries (NumPy, SciPy, librosa) and is used exclusively in this implementation.

Implementation consistency: All implementations must use the same padding type to ensure identical spectrograms for the same input. Using reflective padding when zero-padding is expected (or vice versa) will cause boundary frames to diverge, particularly for the first and last $\lceil N/(2H) \rceil$ frames where padding contributes to the windowed signal.

Implementation detail: The library uses centering by default to avoid edge effects:

Listing 13: Frame count with optional centre-padding

```
fn frame_count(&self, n_samples: usize) -> usize {
    let pad = if self centre { self.n_fft / 2 } else { 0 };
    let padded_len = n_samples + 2 * pad;

    if padded_len < self.n_fft {
```

```

        return 1; // Single frame, even for short signals
    }

    let remaining = padded_len - self.n_fft;
    let n_frames = remaining / self.hop + 1;
    n_frames
}

```

4.4 STFT Computation Pipeline

For each frame m :

1. **Extract frame:** Get samples $x[mH], x[mH + 1], \dots, x[mH + N - 1]$
2. **Apply window:** Compute $x_w[n] = x[mH + n] \cdot w[n]$
3. **FFT:** Compute $X_m[k] = \text{FFT}(x_w)$
4. **Store:** Save as column m of output matrix

Optimizations:

- Pre-compute window samples once
- Reuse FFT plan across all frames
- Allocate workspace buffers once

Workspace Memory Requirements

To enable efficient STFT computation, implementations use a **Workspace** structure that holds reusable buffers and FFT state. This avoids repeated allocations during frame processing.

Minimum workspace requirements:

1. **Windowed frame buffer:** Real-valued array of size N to hold $x[mH + n] \cdot w[n]$ before FFT
2. **FFT output buffer:** Complex-valued array of size $N/2 + 1$ (for R2C FFT) to hold frequency-domain values $X_m[k]$
3. **FFT plan state:** Backend-specific data structure (e.g., FFTW plan, RealFFT planner) that stores twiddle factors and algorithm configuration
4. **Power/magnitude spectrum buffer:** Real-valued array of size $N/2 + 1$ to hold $|X_m[k]|^2$ or $|X_m[k]|$ for the current frame

Total memory requirement: $O(N)$ real values and $O(N)$ complex values, plus backend-specific FFT plan overhead (typically $O(N \log N)$ for precomputed twiddle factors).

Implementation example:

Listing 14: Workspace struct with reusable STFT buffers

```
struct Workspace {
    frame: Vec<f64>,           // Size N (windowed
                               input)
    fft_out: Vec<Complex<f64>>, // Size N/2 + 1 (FFT
                               output)
    spectrum: Vec<f64>,       // Size N/2 + 1 (
                               power/magnitude)
}

impl Workspace {
    fn new(n_fft: usize) -> Self {
        let out_len = n_fft / 2 + 1;
        Self {
            frame: vec![0.0; n_fft],
            fft_out: vec![Complex::zero(); out_len],
            spectrum: vec![0.0; out_len],
        }
    }
}
```

Performance benefit: Reusing workspace buffers across frames reduces memory allocator pressure. For a 10-second audio file at 16 kHz with $N = 512$ and $H = 256$, this processes approximately 625 frames — avoiding 1,875 allocations (3 buffers $\times$ 625 frames).

Listing 15: Windowed frame extraction and FFT per frame

```
- compute one frame
fn compute_frame_spectrum(&mut self, samples: &[f64],
                          frame_idx: usize,
                          workspace: &mut Workspace) ->

Result<()> {
    let out = workspace.frame.as_mut_slice();
    let pad = if self.centre { self.n_fft / 2 } else { 0 };
    let start = frame_idx * self.hop;

    // Fill windowed frame (with zero-padding)
    for i in 0..self.n_fft {
        let v_idx = start + i;
```

```

        let s_idx = v_idx as isize - pad as isize;

        let sample = if s_idx < 0 || s_idx >= samples.len()
        {
            0.0 // Padding
        } else {
            samples[s_idx as usize]
        };

        out[i] = sample * self.window[i];
    }

    // Compute FFT
    self.fft.process(out, workspace.fft_out)?;

    // Convert to power spectrum: |X[k]|^2
    for (i, c) in workspace.fft_out.iter().enumerate() {
        workspace.spectrum[i] = c.norm_sqr();
    }

    Ok(())
}

```

4.5 Spectrogram as STFT Magnitude

A spectrogram is the magnitude (or power) of the STFT:

$$S[m, k] = |X[m, k]|^2 \quad (\text{power spectrogram}) \quad (33)$$

or

$$S[m, k] = |X[m, k]| \quad (\text{magnitude spectrogram}) \quad (34)$$

This discards phase information but provides an intuitive visualization of energy distribution over time and frequency.




Figure 7: Three amplitude representations of the same spectrogram (Hanning window, $N=2048$, 50% hop, $f_s=22,050$ Hz) computed from a four-chord piano progression. **Left:** power $|X[m, k]|^2$ —extremely wide dynamic range makes quiet components invisible. **Centre:** magnitude $|X[m, k]|$ —square-root compression reveals more structure. **Right:** decibel scale $10 \log_{10} |X|^2$ —perceptually linear, with the harmonic series of each chord clearly legible. The dB representation is standard for analysis and display.

4.6 Time-Frequency Uncertainty Principle

Time-Frequency Uncertainty Principle

The fundamental trade-off in STFT is governed by the uncertainty principle [12]:

$$\Delta t \cdot \Delta f \geq \frac{1}{4\pi} \quad (35)$$

where Δt is temporal resolution and Δf is frequency resolution. This is *not* a limitation of the algorithm, but a fundamental property of Fourier analysis related to the Heisenberg uncertainty principle in quantum mechanics. You **cannot** achieve arbitrary precision in both domains simultaneously.

Implications:

- **Large N (long window):** Good frequency resolution, poor time resolution
- **Small N (short window):** Good time resolution, poor frequency resolution

For speech (16 kHz sample rate):

- $N = 512$ (32 ms): $\Delta f = 31.25$ Hz — good for rapid changes
- $N = 2048$ (128 ms): $\Delta f = 7.8$ Hz — better frequency detail

For music (44.1 kHz sample rate):

- $N = 2048$ (46 ms): $\Delta f = 21.5$ Hz — standard choice
- $N = 4096$ (93 ms): $\Delta f = 10.8$ Hz — for harmonic analysis

Cannot achieve arbitrary precision in both domains simultaneously.

This is not a limitation of the algorithm, but a fundamental property of Fourier analysis related to the Heisenberg uncertainty principle in quantum mechanics.




Figure 8: The time-frequency uncertainty trade-off demonstrated on the same piano chord progression at two FFT sizes. **Left** ($N=256$, **31 ms**): individual note onsets are sharply localised in time but the harmonic partials blur together vertically. **Right** ($N=4096$, **185 ms**): individual harmonics are clearly resolved in frequency but chord transitions smear horizontally. Neither window can achieve high resolution in both dimensions simultaneously—this is not an implementation limitation but a fundamental property of Fourier analysis.

5 Amplitude Scaling and Representations

5.1 Power Spectrum

The power spectrum represents energy at each frequency:

$$P[k] = |X[k]|^2 = \text{Re}(X[k])^2 + \text{Im}(X[k])^2 \quad (36)$$

For real input signals, this is proportional to the energy in frequency bin k .

Properties:

- Always non-negative
- Units: (amplitude)²
- Additive: energies from independent sources sum

The squared magnitude emphasizes strong components and suppresses weak ones, which can be both beneficial (noise reduction) and problematic (loss of weak signals).

5.2 Magnitude Spectrum

The magnitude spectrum is the absolute value:

$$M[k] = |X[k]| = \sqrt{\text{Re}(X[k])^2 + \text{Im}(X[k])^2} \quad (37)$$

This is more perceptually linear than power, as human perception of sound intensity is approximately logarithmic.

Relationship:

$$M[k] = \sqrt{P[k]} \quad (38)$$

Use magnitude when:

- You need perceptually meaningful amplitudes
- Comparing with other magnitude-based representations
- Visualization where extreme values would dominate

5.3 Decibel Scale

The decibel (dB) scale provides logarithmic amplitude scaling:

$$S_{\text{dB}}[k] = 10 \log_{10} \left(\frac{P[k]}{P_{\text{ref}}} \right) \quad (39)$$

where P_{ref} is a reference power. For power spectrograms without specific reference:

$$S_{\text{dB}}[k] = 10 \log_{10}(P[k]) \quad (40)$$

For magnitude-based signals:

$$S_{\text{dB}}[k] = 20 \log_{10}(M[k]) \quad (41)$$

The factor of 20 (instead of 10) accounts for the squared relationship between magnitude and power.

Numerical considerations: To avoid $\log(0) = -\infty$, a floor is applied:

$$S_{\text{dB}}[k] = 10 \log_{10}(\max(P[k], \epsilon)) \quad (42)$$

where ϵ is a small threshold (e.g., 10^{-10}). This maps very small values to a large negative dB value instead of $-\infty$.

Dynamic range compression: The dB scale compresses the dynamic range, making it easier to visualize spectrograms with both loud and quiet components. A signal with power ratios of 1,000,000:1 becomes 60 dB, much more manageable.

Example:

- Power = 1.0 $\rightarrow$ 0 dB
- Power = 0.1 $\rightarrow$ -10 dB
- Power = 0.01 $\rightarrow$ -20 dB
- Power = 100.0 $\rightarrow$ 20 dB

5.4 Parseval's Theorem and Energy Conservation

Parseval's Theorem

Parseval's theorem states that energy is conserved between time and frequency domains:

$$\sum_{n=0}^{N-1} |x[n]|^2 = \frac{1}{N} \sum_{k=0}^{N-1} |X[k]|^2 \quad (43)$$

This fundamental result ensures that:

- The FFT doesn't create or destroy energy
- Total power in frequency domain equals total power in time domain
- Inverse FFT perfectly reconstructs energy (up to numerical precision)

For windowed signals in STFT with proper overlap-add, energy is preserved:

$$\sum_m \sum_k |X[m, k]|^2 \propto \sum_n |x[n]|^2 \quad (44)$$

where the proportionality constant depends on the window normalization.

Implication for normalization: Different FFT libraries use different normalization conventions. The library uses the convention where the forward FFT has no normalization factor, and the inverse FFT divides by N . This is consistent with most signal processing literature.




Figure 9: End-to-end feature extraction pipeline applied to synthesised vowel speech (/a/ /e/ /i/ /o/ at $f_s=22\,050$ Hz). **(a)** Raw waveform showing the formant-driven amplitude modulation of each vowel. **(b)** Linear spectrogram (dB, $N=512$): the first two formants (F_1 , F_2) shift noticeably between vowels, encoding their identity. **(c)** Mel spectrogram ($M=40$ bands): perceptually weighted, emphasising the formant region below 3 kHz. **(d)** MFCCs ($C=13$): the DCT decorrelates the mel bands and compacts each vowel’s identity into the first few coefficients—the representation used directly by most speech recognition systems.

6 Perceptual Frequency Scales

6.1 Motivation: Human Auditory Perception

Auditory frequency perception is not linear [9]. Frequency differences are perceived logarithmically, especially at higher frequencies. For example:

- The difference between 100 Hz and 200 Hz is highly perceptible (one octave)
- The difference between 10,000 Hz and 10,100 Hz is barely noticeable (0.014 octaves)

Perceptual frequency scales map linear Hz to psychoacoustically-motivated scales that better match human hearing.

6.2 Mel Scale

The mel scale is based on pitch perception studies [21]. The name comes from “melody,” reflecting its musical origins. The conversion formulas are:

$$\text{mel}(f) = 2595 \log_{10} \left(1 + \frac{f}{700} \right) \quad (45)$$

$$f(\text{mel}) = 700 \left(10^{\text{mel}/2595} - 1 \right) \quad (46)$$

The mel scale is approximately linear below 1000 Hz and logarithmic above.

Mel Filterbank Construction:

To create M mel-spaced filters:

1. Convert frequency range $[f_{\min}, f_{\max}]$ to mel scale
2. Create $M + 2$ equally-spaced mel points (for triangular filter edges)
3. Convert mel points back to Hz
4. Map Hz to FFT bin indices
5. Create triangular filters centered at these points

$$H_m[k] = \begin{cases} 0 & k < f[m-1] \\ \frac{k-f[m-1]}{f[m]-f[m-1]} & f[m-1] \leq k < f[m] \\ \frac{f[m+1]-k}{f[m+1]-f[m]} & f[m] \leq k \leq f[m+1] \\ 0 & k > f[m+1] \end{cases} \quad (47)$$

where $f[m]$ are the FFT bin indices corresponding to mel points.

Implementation detail: The filterbank is stored as a sparse matrix since each triangular filter only has non-zero values over a small range of bins:

Listing 16: Building the sparse mel filterbank matrix

```
fn build_mel_filterbank_matrix(sample_rate_hz: f64, n_fft:
    usize,
                                n_mels: usize, f_min: f64,
                                f_max: f64) -> SparseMatrix
{
    let out_len = r2c_output_size(n_fft);
    let df = sample_rate_hz / n_fft as f64;

    // Convert to mel scale
    let mel_min = hz_to_mel(f_min);
    let mel_max = hz_to_mel(f_max);

    // Create n_mels + 2 mel points (for triangle edges)
    let n_points = n_mels + 2;
    let step = (mel_max - mel_min) / (n_points - 1) as f64;

    let mel_points: Vec<_> = (0..n_points)
        .map(|i| i as f64 * step + mel_min)
        .collect();

    // Convert back to Hz, then to FFT bins
    let bin_points: Vec<_> = mel_points
        .iter()
        .map(|&mel| mel_to_hz(mel))
        .map(|hz| (hz / df).floor() as usize)
        .collect();

    // Build sparse filterbank
    let mut fb = SparseMatrix::new(n_mels, out_len);

    for m in 0..n_mels {
        let left = bin_points[m];
        let centre = bin_points[m + 1];
        let right = bin_points[m + 2];

        // Rising slope: left -> centre
        for k in left..centre {
            let v = (k - left) as f64 / (centre - left) as
f64;
            fb.set(m, k, v);
        }
    }
}
```

```

    // Falling slope: centre -> right
    for k in centre..right {
        let v = (right - k) as f64 / (right - centre) as
f64;
        fb.set(m, k, v);
    }
}

fb
}

```

For typical speech processing ($N = 512$, 80 mel bands), the filterbank is approximately 95% sparse.

Numerical hygiene: To achieve this sparsity in practice, coefficients below a threshold (typically 10^{-12} or machine epsilon $\epsilon_{\text{machine}} \approx 2.2 \times 10^{-16}$ for f64) are set to exactly zero during construction. This prevents accumulation of negligible numerical artifacts while maintaining filterbank fidelity.




Figure 10: **Top:** mel triangular filterbank ($M=40$ filters, $f_{\min}=0$ Hz, $f_{\max}=8$ kHz). Each overlapping triangle is wider at higher frequencies (more Hz per mel), capturing the logarithmic perception of pitch. Colour indicates filter index from low (violet) to high (yellow). **Bottom:** the Hz-to-mel mapping $\text{mel}(f) = 2595 \log_{10}(1 + f/700)$, linear below 1 kHz and logarithmic above.




Figure 11: Mel-scale spectrogram (dB) of a four-chord piano progression (C maj → A min → F maj → G maj), $M=128$ mel bands. The perceptual frequency axis allocates many bands to the musically rich low-frequency region and fewer to the upper overtones, matching the logarithmic sensitivity of human pitch perception. The harmonic series of each chord is clearly visible as parallel horizontal streaks.

6.3 ERB Scale: Glasberg and Moore Model

The *Equivalent Rectangular Bandwidth* (ERB) scale is based on psychoacoustic measurements of critical bandwidths [14, 17]. It models the frequency selectivity of the human auditory system.

The ERB in Hz as a function of center frequency:

$$\text{ERB}(f) = 24.7 \left(4.37 \frac{f}{1000} + 1 \right) \quad (48)$$

To convert frequency to ERB scale (ERB-rate):

$$\text{ERB-rate}(f) = 21.4 \log_{10} \left(4.37 \frac{f}{1000} + 1 \right) \quad (49)$$

Inverse transformation:

$$f(\text{ERB}) = \frac{1000}{4.37} \left(10^{\text{ERB}/21.4} - 1 \right) \quad (50)$$

Gammatone Filters:

ERB spectrograms often use gammatone filters, which model the impulse response of the auditory filter:

$$g(t) = t^{n-1} e^{-2\pi b t} \cos(2\pi f_c t + \phi) \quad (51)$$

where:

- n is the filter order (typically 4)
- b is the bandwidth parameter related to ERB
- f_c is the center frequency
- ϕ is the phase

The library implements gammatone filters in the frequency domain for efficiency.

Why ERB over Mel? The ERB scale more accurately models auditory filters at higher frequencies and is preferred for applications requiring high fidelity to human perception.




Figure 12: Gammatone auditory filters at five centre frequencies (125 Hz, 500 Hz, 1 kHz, 2 kHz, 4 kHz). **Left:** time-domain impulse responses (order $n=4$, bandwidth $b=1.019 \cdot \text{ERB}(f_c)$). The envelope $t^{n-1}e^{-2\pi b t}$ shortens at higher frequencies because the ERB bandwidth grows with f_c . **Right:** frequency responses showing the characteristic asymmetric skirt of the auditory filter—sharper on the high-frequency side, broader toward lower frequencies.

6.4 Logarithmic Frequency (Musical Scale)

Musical notes are logarithmically spaced: each octave doubles the frequency. For music analysis, a logarithmic frequency axis is natural.

Log-frequency bins:

$$f_k = f_{\min} \cdot \left(\frac{f_{\max}}{f_{\min}} \right)^{k/(K-1)} \quad (52)$$

for $k = 0, 1, \dots, K-1$ bins.

Alternatively, using equal spacing in log-space:

$$\log f_k = \log f_{\min} + k \frac{\log f_{\max} - \log f_{\min}}{K-1} \quad (53)$$

Linear Interpolation:

Since FFT bins are linearly spaced, linear interpolation is used between adjacent bins to get log-spaced values:

$$y_k = (1 - \alpha)X[i] + \alpha X[i + 1] \quad (54)$$

where $i = \lfloor f_k / \Delta f \rfloor$ and $\alpha = (f_k / \Delta f) - i$.

Listing 17: Log-frequency interpolation via sparse matrix

```
fn build_loghz_matrix(sample_rate_hz: f64, n_fft: usize,
                      n_bins: usize, f_min: f64,
                      f_max: f64) -> SparseMatrix {
    let df = sample_rate_hz / n_fft as f64;

    // Logarithmically-spaced frequencies
    let log_f_min = f_min.ln();
    let log_f_max = f_max.ln();
    let log_step = (log_f_max - log_f_min) / (n_bins - 1) as
        f64;

    let log_frequencies: Vec<_> = (0..n_bins)
        .map(|i| (i as f64 * log_step + log_f_min).exp())
        .collect();

    // Build interpolation matrix
    let mut matrix = SparseMatrix::new(n_bins, out_len);

    for (bin_idx, &target_freq) in log_frequencies.iter().
        enumerate() {
        let exact_bin = target_freq / df;
        let lower_bin = exact_bin.floor() as usize;
        let upper_bin = exact_bin.ceil() as usize;

        if lower_bin == upper_bin {
            matrix.set(bin_idx, lower_bin, 1.0);
        } else {
            // Linear interpolation
            let frac = exact_bin - lower_bin as f64;
            matrix.set(bin_idx, lower_bin, 1.0 - frac);
            matrix.set(bin_idx, upper_bin, frac);
        }
    }

    matrix
}
```


This matrix is extremely sparse: only 1–2 non-zero values per row (99% sparse).

6.5 Filterbank Design and Sparse Matrices

All perceptual scales use sparse matrix multiplication for efficiency. The operation is:

$$\mathbf{y} = H\mathbf{x} \quad (55)$$

where H is the filterbank matrix (sparse), $\mathbf{x}$ is the linear spectrum, and $\mathbf{y}$ is the perceptually-scaled spectrum.

Sparse Matrix Storage:

The library uses row-wise sparse format:

Listing 18: Row-wise sparse matrix storage definition

```
struct SparseMatrix {
    nrows: usize,
    ncols: usize,
    values: Vec<Vec<f64>>, // Non-zero values per row
    indices: Vec<Vec<usize>>, // Column indices per row
}
```

Multiplication:

Listing 19: Sparse matrix-vector multiplication

```
fn multiply_vec(&self, input: &[f64], out: &mut [f64]) {
    for (row_idx, (row_values, row_indices)) in
        self.values.iter().zip(&self.indices).enumerate() {

        let mut acc = 0.0;
        for (&value, &col_idx) in row_values.iter().zip(
            row_indices) {
            acc += value * input[col_idx];
        }
        out[row_idx] = acc;
    }
}
```

Sparse Matrix Computational Complexity

For matrix-vector multiplication $\mathbf{y} = A\mathbf{x}$ where A is $m \times n$:

Dense matrix:

$$\mathcal{O}(mn) \quad (\text{every element accessed}) \quad (56)$$

Sparse matrix with nnz non-zeros:

$$\mathcal{O}(\text{nnz}) \quad (\text{only non-zeros computed}) \quad (57)$$

where nnz is the number of non-zero elements.

For filterbanks with uniform sparsity s non-zeros per row:

$$\text{nnz} = ms \quad (58)$$

Example: Mel filterbank with $m = 80$, $n = 257$, $s = 20$ non-zeros/row has 1,600 operations vs. 20,560 for dense multiplication.

For mel filterbanks with $M = 80$ bands and $N = 512$ FFT (257 bins):

- Dense matrix: $80 \times 257 = 20,560$ multiplications
- Sparse matrix: $\approx 80 \times 20 = 1,600$ multiplications (~ 10 – 20 non-zeros/row)

For log-frequency with linear interpolation:

- Only 2 multiplications per output bin (vs. 257 for dense)

7 Constant-Q Transform (CQT)

7.1 Motivation and Q Factor

The Short-Time Fourier Transform (Section 4.2) provides constant frequency resolution Δf across all frequencies. But musical perception is logarithmic: *relative* frequency differences (ratios) are more relevant than absolute differences.

The **Constant-Q Transform** (CQT) [5, 6] provides logarithmically-spaced frequency bins with constant quality factor Q :

$$Q = \frac{f_k}{\Delta f_k} \quad (59)$$

where f_k is the center frequency and Δf_k is the bandwidth of bin k .

For musical applications with B bins per octave:

$$Q = \frac{1}{2^{1/B} - 1} \quad (60)$$

For example, with $B = 12$ bins/octave (semitones):

$$Q = \frac{1}{2^{1/12} - 1} \approx 16.82 \quad (61)$$

Higher $Q \rightarrow$ narrower bandwidth $\rightarrow$ better frequency resolution but worse time resolution.

7.2 Kernel-Based Implementation

Unlike STFT which uses a fixed-length window, CQT uses *variable-length windows* (kernels) for each frequency:

For frequency bin k with center frequency f_k :

$$N_k = \left\lceil \frac{Q \cdot f_s}{f_k} \right\rceil \quad (62)$$

where f_s is the sample rate. Higher frequencies $\rightarrow$ shorter windows, lower frequencies $\rightarrow$ longer windows.

The CQT kernel for frequency f_k :

$$g_k[n] = w[n] \cdot e^{j2\pi f_k n / f_s}, \quad n = 0, 1, \dots, N_k - 1 \quad (63)$$

where $w[n]$ is a window function (typically Hanning or Hamming).

Implementation:

Listing 20: CQT kernel with windowed complex exponential

```
fn generate_kernel_bin(center_freq: f64, kernel_length:
    usize,
                        sample_rate: f64,
                        window_type: WindowType) -> Vec<
Complex<f64>> {
    let mut kernel = Vec::with_capacity(kernel_length);

    // Generate window
    let window = make_window(window_type, kernel_length)?;
```

```

// Generate complex exponential kernel
for (n, w) in window.iter().enumerate() {
    let t = n as f64 / sample_rate;
    let phase = 2.0 * PI * center_freq * t;

    //  $e^{i2\pi f t} = \cos(2\pi f t) + i\sin(2\pi f t)$ 
    let exponential = Complex::new(phase.cos(), phase.sin());

    // Apply window function
    let windowed = exponential * w;
    kernel.push(windowed);
}

kernel
}

```

7.3 Variable-Length Windows

The beauty of CQT is that window length scales with frequency:

- **Low frequencies:** Long windows $\rightarrow$ good frequency resolution, poor time resolution
- **High frequencies:** Short windows $\rightarrow$ good time resolution, adequate frequency resolution

This matches musical perception: precise pitch resolution is required for low notes (fundamental frequencies), while less precision is typically acceptable for high harmonics.

Example for $f_s = 44,100$ Hz, $Q \approx 16.82$:

- $f = 55$ Hz (A1): $N \approx 13,480$ samples (≈ 305 ms)
- $f = 440$ Hz (A4): $N \approx 1,685$ samples (≈ 38 ms)
- $f = 3,520$ Hz (A7): $N \approx 211$ samples (≈ 4.8 ms)

7.4 Sparsity Optimization

CQT kernels naturally have many small-magnitude coefficients, especially near the edges due to windowing. This can be exploited with *sparsity thresholding*:

$$g_k[n] \leftarrow \begin{cases} g_k[n] & \text{if } |g_k[n]| > \epsilon \cdot \max |g_k| \\ 0 & \text{otherwise} \end{cases} \quad (64)$$

where ϵ is the sparsity threshold (e.g., 0.01 for 1%).

Listing 21: Zeroing small CQT kernel coefficients

```
fn apply_sparsity_threshold(kernel: &mut [Complex<f64>],
                           threshold: f64) {
    let max_magnitude = kernel.iter()
        .map(|c| c.norm())
        .fold(0.0, f64::max);

    let absolute_threshold = max_magnitude * threshold;

    for coefficient in kernel.iter_mut() {
        if coefficient.norm() < absolute_threshold {
            *coefficient = Complex::new(0.0, 0.0);
        }
    }
}
```

Typical sparsity levels: 60–80% of kernel coefficients can be zeroed with $\epsilon = 0.01$ with negligible quality loss.

Computational benefit: Sparse kernels reduce the number of operations in time-domain convolution by skipping zero multiplications.

Sparsity Threshold Artifacts

Warning: Aggressive sparsity thresholds (e.g., $\epsilon > 0.01$ or zeroing coefficients below 10^{-10}) can introduce audible artifacts:

- **Spectral leakage:** Truncating kernel tails creates discontinuities, causing frequency bleeding between adjacent bins.
- **Musical noise:** Sparse, time-varying kernels can produce tonal artifacts that vary frame-to-frame, perceived as “chirping” or “bubbling” sounds.

- **Time-smearing:** Excessively truncated kernels lose their designed frequency selectivity.

Recommended practice: Use conservative thresholds ($\epsilon \leq 0.001$) for high-quality audio analysis. For real-time or resource-constrained applications, validate perceptual quality through listening tests before deploying aggressive sparsity.




Figure 13: Constant-Q Transform spectrogram of a chromatic ascending arpeggio from C2 to C8 (12 bins/octave, 6 octaves, $f_{\min}=65$ Hz, dB scale). Unlike the STFT, successive semitones are uniformly spaced vertically—a direct consequence of the constant- Q design. Each note produces a brief burst that aligns precisely with its pitch-class label on the y -axis, demonstrating the transform’s suitability for pitch-accurate music analysis.

8 Advanced Audio Features

8.1 Mel-Frequency Cepstral Coefficients (MFCC)

MFCCs are the most widely used features in speech recognition, speaker identification, and audio classification [8, 20]. Introduced by Davis and Mermelstein in 1980, they provide a compact representation of the spectral envelope.

Computation pipeline:

1. Compute mel-scaled power spectrogram: $S_{\text{mel}}[m, k]$
2. Apply logarithm: $L[m, k] = \log(S_{\text{mel}}[m, k] + \epsilon)$
3. Apply Discrete Cosine Transform (DCT-II) to each frame
4. Optionally apply cepstral liftering
5. Extract first C coefficients (typically 13)

Why DCT? The mel spectrogram has correlated bins due to overlapping triangular filters. The *Discrete Cosine Transform* (DCT) [1] decorrelates them and compacts energy into the first few coefficients.

Discrete Cosine Transform (DCT-II)

The DCT-II formula:

$$c[k] = \sum_{n=0}^{N-1} x[n] \cos\left(\frac{\pi k(n+0.5)}{N}\right), \quad k = 0, 1, \dots, N-1 \quad (65)$$

The DCT is closely related to the DFT but operates entirely on real numbers. It's used in JPEG image compression and many audio codecs due to its excellent energy compaction properties.

Listing 22: DCT-II implementation for MFCC cepstral step

```
fn dct_ii(input: &[f64]) -> Vec<f64> {
    let n = input.len();
    let mut output = vec![0.0; n];

    for k in 0..n {
        let mut sum = 0.0;
        for (i, &val) in input.iter().enumerate() {
            sum += val * (PI * k as f64 * (i as f64 + 0.5) /
n as f64).cos();
        }
        output[k] = sum;
    }
}
```

```
    output
}
```

Cepstral liftering:

Liftering applies a sinusoidal weighting to emphasize mid-range coefficients:

$$\tilde{c}[k] = c[k] \left(1 + \frac{L}{2} \sin \left(\frac{\pi k}{L} \right) \right) \quad (66)$$

where L is the lifter parameter (commonly 22).

Listing 23: Sinusoidal liftering applied to MFCC array

```
fn apply_liftering(mfcc: &mut Array2<f64>, lifter: usize) {
    let n_mfcc = mfcc.nrows();

    // Compute lifter weights
    let weights: Vec<_> = (0..n_mfcc)
        .map(|i| 1.0 + (lifter as f64 / 2.0)
            * (PI * i as f64 / lifter as f64).sin())
    )
        .collect();

    // Apply weights to each frame
    for frame_idx in 0..n_frames {
        for coeff_idx in 0..n_mfcc {
            mfcc[[coeff_idx, frame_idx]] *= weights[
                coeff_idx];
        }
    }
}
```

The C0 coefficient: The zeroth coefficient represents the overall energy of the frame. Some applications include it (speaker recognition), others exclude it (speech recognition) to reduce sensitivity to volume.

Why 13 coefficients? Empirically, 13 MFCCs capture most speech information. Higher coefficients represent fine spectral detail that's often noise.




Figure 14: MFCCs ($C=13$) computed from five synthesised vowels (/a/, /e/, /i/, /o/, /u/, each ≈ 0.55 s). Vertical white dashes mark vowel transitions. The low-order coefficients (rows 0–4) encode the broad spectral envelope (vowel identity), while higher-order coefficients capture finer spectral detail. The DCT-II decorrelates the mel-band energies, concentrating the discriminative information into the first few coefficients—which is why 13 MFCCs are sufficient for most speech tasks.

8.2 Chroma Features for Music Analysis

Chroma features [11, 18] (also called pitch class profiles) represent the energy distribution across the 12 pitch classes of the Western musical scale, collapsing all octaves together.

The 12 pitch classes: C, C $\sharp$, D, D $\sharp$, E, F, F $\sharp$, G, G $\sharp$, A, A $\sharp$, B

Construction:

1. Map FFT bins to musical notes using:

$$n(f) = 12 \log_2 \left(\frac{f}{f_{\text{ref}}} \right) \quad (67)$$

where f_{ref} is a reference frequency (e.g., A4 = 440 Hz)

2. Determine pitch class: $p = n \bmod 12$

3. Sum energy from all octaves for each pitch class:

$$C[p] = \sum_{\text{octaves}} E[p, \text{octave}] \quad (68)$$

4. Normalize (L1, L2, or max)

Normalization strategies:

- **L1:** $\tilde{C}[p] = C[p] / \sum_{p'} C[p']$ (probabilities)
- **L2:** $\tilde{C}[p] = C[p] / \sqrt{\sum_{p'} C[p']^2}$ (Euclidean norm)
- **Max:** $\tilde{C}[p] = C[p] / \max_{p'} C[p']$ (relative strength)

Applications: Chord recognition, key estimation, cover song identification, music similarity.




Figure 15: Chromagram of the four-chord progression (C maj → A min → F maj → G maj). The y -axis lists the 12 pitch classes; energy is summed across all octaves, so the display is octave-invariant. Within the C major segment (0–1 s), the C, E, and G rows light up simultaneously; A minor activates A, C, and E; and so on. This compact 12-dimensional vector is sufficient for key detection, chord recognition, and cover-song identification.

8.3 ERB Spectrograms and Gammatone Filters

ERB spectrograms using gammatone filters provide the most perceptually accurate time-frequency representation.

The gammatone impulse response:

$$g(t) = at^{n-1}e^{-2\pi bt}\cos(2\pi f_c t + \phi) \quad (69)$$

where:

- a is amplitude
- n is filter order (typically 4)
- b is bandwidth parameter
- f_c is center frequency
- ϕ is phase

The bandwidth parameter relates to ERB:

$$b = 1.019 \cdot \text{ERB}(f_c) \quad (70)$$

Frequency domain implementation:

For efficiency, the library implements gammatone filtering in the frequency domain. The frequency-domain transfer function is:

$$G(f) = \frac{1}{\left(1 + j\frac{f-f_c}{b}\right)^n} \quad (71)$$

where n is the filter order (typically 4), f_c is the center frequency, b is the bandwidth parameter, and $j = \sqrt{-1}$. This fourth-order complex pole provides the characteristic asymmetric frequency response of the auditory filter.

Filtering is then performed via element-wise multiplication:

$$Y(f) = X(f) \cdot G(f) \quad (72)$$

This avoids expensive time-domain convolution for each filter.




Figure 16: ERB/gammatone spectrogram of the chord progression (log-power scale, 64 channels from 80 Hz to 8 kHz). Each channel’s bandwidth grows with frequency following the ERB function, providing finer frequency resolution in the low-frequency region where harmonic relationships are most salient.

9 Image Processing via FFT

9.1 2D Convolution and the Convolution Theorem

The *convolution theorem* [4] states that convolution in the spatial domain is equivalent to multiplication in the frequency domain:

$$f * g = \mathcal{F}^{-1}\{\mathcal{F}\{f\} \cdot \mathcal{F}\{g\}\} \quad (73)$$

For images:

$$(I * K)[x, y] = \text{IFFT2D}\{\text{FFT2D}(I) \odot \text{FFT2D}(K)\} \quad (74)$$

where $\odot$ denotes element-wise multiplication.

Complexity comparison:

Spatial convolution for image $M \times N$ with kernel $K \times K$:

$$\mathcal{O}(M \times N \times K^2) \quad (75)$$

FFT-based convolution:

$$\mathcal{O}(MN \log(MN)) \quad (76)$$

Theoretical crossover: FFT has better asymptotic complexity when $K > \sqrt{\log(MN)}$, typically around $K \geq 7$.

9.2 Kernel Padding and Phase Shifting

For correct FFT convolution, the kernel must be padded to image size with proper phase shifting. The kernel's center should map to position (0,0) in the padded array:

Listing 24: Padding kernel with centre at origin for FFT

```
fn pad_kernel_for_fft(kernel: &Array2<f64>,
                      target_shape: (usize, usize)) ->
    Array2<f64> {
    let (ker_rows, ker_cols) = kernel.dim();
    let mut result = Array2::::zeros(target_shape);

    // Kernel center position
    let ker_center_row = ker_rows / 2;
    let ker_center_col = ker_cols / 2;

    // Place kernel with center at (0, 0) using wraparound
    for i in 0..ker_rows {
        for j in 0..ker_cols {
            let row_offset = i as isize - ker_center_row as
isize;
            let col_offset = j as isize - ker_center_col as
isize;

            // Wrap around to place center at (0, 0)
            let target_row = row_offset.rem_euclid(
target_rows as isize) as usize;
            let target_col = col_offset.rem_euclid(
target_cols as isize) as usize;

            result[[target_row, target_col]] = kernel[[i, j
]];
        }
    }
```

```

    }

    result
}

```

This ensures the FFT interprets the kernel correctly for circular convolution.

9.3 Spatial Filtering in the Frequency Domain

Filtering is multiplication by a frequency-domain mask:

$$Y(u, v) = X(u, v) \cdot H(u, v) \quad (77)$$

Low-pass filter: Keeps low frequencies, removes high frequencies (smoothing)

$$H_{\text{LP}}(u, v) = \begin{cases} 1 & \sqrt{u^2 + v^2} \leq D_0 \\ 0 & \text{otherwise} \end{cases} \quad (78)$$

High-pass filter: Removes low frequencies, keeps high frequencies (sharpening/edges)

$$H_{\text{HP}}(u, v) = 1 - H_{\text{LP}}(u, v) \quad (79)$$

Band-pass filter: Keeps frequencies in a range

$$H_{\text{BP}}(u, v) = \begin{cases} 1 & D_1 \leq \sqrt{u^2 + v^2} \leq D_2 \\ 0 & \text{otherwise} \end{cases} \quad (80)$$

9.4 Gaussian Blur via FFT

The Gaussian kernel for blurring:

$$G(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \quad (81)$$

Implementation:

Listing 25: 2D Gaussian kernel generation and normalisation

```

pub fn gaussian_kernel_2d(size: NonZeroUsize, sigma: f64) ->
    SpectrogramResult<Array2<f64>> {
    if size.get().is_multiple_of(2) {
        return Err(SpectrogramError::invalid_input(
            "kernel size must be odd and > 0",
        ));
    }
    if sigma <= 0.0 {
        return Err(SpectrogramError::invalid_input("sigma
must be > 0"));
    }
    let size = size.get();
    let center = (size / 2) as f64;
    let variance = sigma * sigma;
    let coeff = 1.0 / (2.0 * PI * variance);

    let mut kernel = Array2::<f64>::zeros((size, size));

    for i in 0..size {
        for j in 0..size {
            let x = i as f64 - center;
            let y = j as f64 - center;
            let exponent = -(x * x + y * y) / (2.0 *
variance);
            kernel[[i, j]] = coeff * exponent.exp();
        }
    }

    // Normalize to sum to 1.0
    let sum: f64 = kernel.iter().sum();
    kernel.mapv_inplace(|v| v / sum);

    Ok(kernel)
}

```

The Gaussian has a useful property: its Fourier transform is also Gaussian, ensuring smooth frequency response.

9.5 Edge Detection and High-Pass Filtering

Edges correspond to high-frequency components. A simple high-pass filter enhances edges:

$$\text{Edges} = \text{Image} - \text{Lowpass}(\text{Image}) \quad (82)$$

Or apply a high-pass mask directly in frequency domain:

$$H_{\text{HP}}(u, v) = \begin{cases} 0 & \sqrt{u^2 + v^2} < D_0 \\ 1 & \text{otherwise} \end{cases} \quad (83)$$

The Laplacian operator for edge detection:

$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2} \quad (84)$$

In frequency domain, differentiation becomes multiplication:

$$\mathcal{F}\{\nabla^2 I\} = -(u^2 + v^2)\mathcal{F}\{I\} \quad (85)$$




Figure 17: Image processing pipeline via 2D FFT applied to a 320×320 greyscale crop of the raccoon portrait (`scipy.datasets.face`). **(1) Original:** the unmodified input. **(2) 2D FFT magnitude ($\log |F|$):** the DC component (centre) is the brightest point; horizontal and vertical lines correspond to strong directional content in the image. **(3) Gaussian blur ($\sigma=4$):** low-pass filtering via multiplication by a Gaussian mask in the frequency domain. **(4) High-pass edges ($r < 20$):** zeroing the central disc retains only high-frequency content, revealing structural edges.

10 Binaural Spectrograms

10.1 Spatial Hearing and Binaural Cues

The human auditory system localises sound by comparing the signals arriving at the two ears. The library exposes four classic binaural cues as time-frequency representations, each implemented as a spectrogram whose rows correspond to frequency bins and whose columns correspond to STFT frames. All four operate at either f32 or f64 precision.

Cue	Symbol	Units	Frequency range
Interaural Time Difference	ITD	seconds	<1.5 kHz
Interaural Phase Difference	IPD	radians	all
Interaural Level Difference	ILD	dB	>1.5 kHz
Interaural Level Ratio	ILR	dimensionless $([-1, 1])$	all

10.2 Interaural Time Difference (ITD)

ITD Definition

The ITD is the extra propagation time for a sound to reach the farther ear:

$$\text{ITD}(k, m) = \frac{\angle X_L[k, m] - \angle X_R[k, m]}{2\pi f_k} \quad (86)$$

where X_L and X_R are the complex STFT spectra of the left and right channels, $\angle$ denotes argument (phase angle), and f_k is the centre frequency of bin k in Hz. At low frequencies the phase difference is an unambiguous indicator of ITD; above ~ 1.5 kHz phase wrapping makes the relationship ambiguous.

The library uses `round` (not `floor/ceil`) for the bin-index conversion to match the `Binaspect` reference implementation, and returns only the bins spanning `[start_freq, end_freq]` rather than the full spectrum:

Listing 26: Computing ITD spectrogram from stereo channels

```
use spectrograms::{ITDSpectrogramParams,
    compute_itd_spectrogram};
use std::num::NonZeroUsize;

let n_fft = NonZeroUsize::new(4096).unwrap();
let hop = NonZeroUsize::new(2048).unwrap();
```

```

let params = ITDSpectrogramParams::new(50.0, 620.0, n_fft,
    hop)?;
// start_freq=50 Hz, end_freq=620 Hz

let itd = compute_itd_spectrogram(&left, &right, sample_rate,
    &params)?;
// itd is an ItdSpectrogram<f64>; shape is (n_freq_bins,
    n_frames)

```

10.3 Interaural Phase Difference (IPD)

The IPD is the unwrapped phase difference between channels:

$$\text{IPD}(k, m) = \angle X_L[k, m] - \angle X_R[k, m] \quad (87)$$

Values lie in $(-\pi, \pi]$. Unlike ITD, IPD is frequency-independent and does not require dividing by f_k , making it ambiguous at high frequencies (half a wavelength becomes less than the inter-ear distance).

10.4 Interaural Level Difference (ILD)

ILD Definition

The ILD captures level asymmetry between channels in decibels:

$$\text{ILD}(k, m) = 10 \log_{10} \frac{|X_L[k, m]|^2 + \epsilon}{|X_R[k, m]|^2 + \epsilon} \quad (88)$$

where ϵ prevents division by zero. Positive values indicate the sound source is closer to the left ear. ILD is the dominant cue above ~ 1.5 kHz because the head creates a significant acoustic shadow at those wavelengths.

10.5 Interaural Level Ratio (ILR)

The ILR normalises the power difference to the $[-1, 1]$ range:

$$\text{ILR}(k, m) = \frac{|X_L[k, m]|^2 - |X_R[k, m]|^2}{|X_L[k, m]|^2 + |X_R[k, m]|^2 + \epsilon} \quad (89)$$

A value of +1 indicates all energy in the left channel; -1 indicates all energy in the right channel; 0 is perfectly balanced.

10.6 Histogram Analysis and Comparison

Every binaural spectrogram object exposes a `histogram(n.bins)` method that aggregates the frame-by-frame cue values into a 2D histogram (frequency bins $\times$ value buckets), providing a statistical summary of the spatial distribution over time.

The library also provides comparison functions that compute the mean per-frequency-bin difference between a reference and a test recording:

Listing 27: Mean ITD difference between reference and test

```
let (col_means, mean_diff_degrees, mean_diff_itd) =  
  compute_itd_spectrogram_diff(&ref_left, &ref_right,  
                                &test_left, &test_right,  
                                sample_rate, &params)?;  
  
// col_means           -- per-frequency-bin mean ITD  
//                      difference (Array1)  
// mean_diff_degrees -- mean absolute error converted to  
//                      degrees of azimuth  
// mean_diff_itd      -- median ITD difference across frames  
//                      (seconds)
```

10.7 Parallelisation

When the `rayon` feature is enabled, the per-frequency-bin magphase computation uses `par_for_each`, giving a significant speed-up on multi-core systems. Enable it in `Cargo.toml`:

Listing 28: Enabling rayon parallelisation in `Cargo.toml`

```
[dependencies]
spectrograms = { version = "2", features = ["rayon"] }
```




Figure 18: Binaural spectrograms for a simulated two-source scene: a piano chord at 30° right ($\text{ITD} \approx 340 \mu\text{s}$, $\text{ILD} \approx 8 \text{ dB}$) mixed with a chirp at 60° left. All four cues are shown as time-frequency maps (red = rightward/positive, blue = leftward/negative). **ITD** is most reliable at low frequencies ($< 1.5 \text{ kHz}$); **IPD** wraps at high frequencies due to sub-wavelength ambiguity; **ILD** dominates above 1.5 kHz where the head creates an acoustic shadow; **ILR** normalises the power asymmetry to $[-1, +1]$ and is interpretable at all frequencies.

11 Modified Discrete Cosine Transform (MDCT)

11.1 Motivation and Codec Context

The *Modified Discrete Cosine Transform* (MDCT) is a lapped orthogonal transform at the heart of virtually every modern lossy audio codec: MP3 (Layer III), AAC, Vorbis, and Opus all use it as their core time-frequency analysis stage. Unlike the STFT, the MDCT overlaps input frames by exactly 50% and maps $2N$ samples to N real-valued coefficients (50% dimensionality reduction), yet still achieves *perfect reconstruction* when the appropriate window is applied.

11.2 MDCT Definition

For a windowed frame of $2N$ samples $x[n]$, the MDCT produces N coefficients:

$$C[k] = \sum_{n=0}^{2N-1} x[n] w[n] \cos\left(\frac{\pi(2n+1+N)(2k+1)}{4N}\right), \quad k = 0, 1, \dots, N-1 \quad (90)$$

where $w[n]$ is a window function satisfying the TDAC condition (see Section 11.4).

The inverse MDCT (IMDCT) reconstructs $2N$ samples from N coefficients:

$$y[n] = \frac{2}{N} \sum_{k=0}^{N-1} C[k] w[n] \cos\left(\frac{\pi(2n+1+N)(2k+1)}{4N}\right) \quad (91)$$

Consecutive frames overlap by N samples and are summed (*overlap-add*); the TDAC window ensures cancellation of aliasing terms.

11.3 Fast Computation via the Packing Trick

Direct evaluation of Equation (90) costs $O(N^2)$. The library computes the forward MDCT in $O(N \log N)$ using a single N -point complex-to-complex (C2c) FFT (the *packing trick*):

1. Pre-multiply each windowed sample by a twiddle factor to produce a complex sequence $z[m] = a[m] + jb[m]$ of length N from the $2N$ -sample frame.

2. Execute one N -point C2c FFT on z .
3. Because a and b are real, the result satisfies Hermitian symmetry; the $2N$ -point DFT at even and odd indices is recovered via:

$$F[2p] = Z[p], \quad F[2p + 1] = Z[N - 1 - p]^* \quad (92)$$

4. Apply a post-twiddle to obtain the MDCT coefficients.

This halves the FFT count compared to the naive $2 \times \text{R2c}(N)$ approach.

The inverse (IMDCT) similarly uses $2 \times \text{C2c}(N)$ FFTs by splitting the half-sparse $2N$ -point DFT into even- and odd-indexed subsequences.

11.4 Perfect Reconstruction and the TDAC Condition

Time-Domain Aliasing Cancellation (TDAC)

Perfect reconstruction requires the window w to satisfy:

$$w[n]^2 + w[n + N]^2 = 1, \quad \forall n \in [0, N) \quad (93)$$

Only the sine window satisfies this condition for 50% hop:

$$w[n] = \sin\left(\frac{\pi(n + \frac{1}{2})}{2N}\right) \quad (94)$$

Standard windows (Hanning, Hamming, Blackman, Kaiser) do *not* satisfy Equation (93) and will introduce reconstruction error.

Usage:

Listing 29: MDCT and IMDCT with sine-window reconstruction

```
use spectrograms::{MdctParams, mdct, imdct};
use std::num::NonZeroUsize;

// Sine window gives perfect reconstruction with 50% hop.
let params = MdctParams::sine_window(NonZeroUsize::new(2048)
    .unwrap())?;

let frames: Vec<Vec<f64>> = /* overlapping 2048-sample
    frames */;
let coeffs = mdct(&frames, &params)?; // (n_frames,
    1024) coefficients
```

```
let reconstructed = imdct(&coeffs, &params)?; // overlap-
      added output
```

Single precision:

Listing 30: MDCT at single-precision f32 via turbofish

```
let coeffs = mdct::<f32>(&frames_f32, &params)?;
```

11.5 Parameters

Parameter	Type	Notes
window_size	NonZeroUsize (even, ≥ 4)	Total frame length $2N$
hop_size	NonZeroUsize	Typically N for PR
window	WindowType	Use <code>MdctParams::sine_window</code> for PR




Figure 19: MDCT analysis and perfect reconstruction of a 2-second piano chord progression (sine window, $2N=2048$, 50% hop). **Top:** MDCT coefficient matrix $\log_{10} |C[k, m]|$, analogous to a spectrogram but with 50% dimensionality reduction: N input samples map to $N/2$ real coefficients per frame. **Centre:** original waveform (blue) overlaid with the IMDCT reconstruction (red dashed)—the two are visually indistinguishable. **Bottom:** reconstruction error, which is at the level of floating-point rounding noise ($\approx 10^{-14}$), confirming the Time-Domain Aliasing Cancellation (TDAC) property of the sine window.

12 FFT Convolution and Minimum Phase

12.1 FFT-Based Linear Convolution

The convolution theorem (Section 9.1) applies equally to 1D real signals. For two signals a of length M and b of length K , linear convolution yields an output of length $M + K - 1$. Direct time-domain convolution costs $O(MK)$; the FFT-based approach costs $O((M + K) \log(M + K))$ and is more efficient whenever $K \gtrsim \sqrt{\log(M + K)}$.

Listing 31: FFT linear convolution with NonEmptySlice input

```
use spectrograms::fft_convolve;
use non_empty_slice::NonEmptySlice;

let signal: Vec<f64> = /* ... */;
let kernel: Vec<f64> = /* FIR filter taps */;

let sig_ne = NonEmptySlice::new(&signal).unwrap();
let ker_ne = NonEmptySlice::new(&kernel).unwrap();

let output = fft_convolve(sig_ne, ker_ne)?;
// output.len() == signal.len() + kernel.len() - 1
```

Algorithm: Both inputs are zero-padded to the next power-of-two above $M + K - 1$, multiplied element-wise in the frequency domain (a single complex multiplication per frequency bin), and inverse-transformed.

12.2 Spectral-Division Deconvolution

Given an observed output $y = h * x$ and the filter h , the original signal x can be recovered by dividing spectra. Exact spectral division is unstable at frequencies where $H(f) \approx 0$; **regularised deconvolution** addresses this:

$$\hat{X}(f) = \frac{Y(f) H^*(f)}{|H(f)|^2 + \varepsilon} \quad (95)$$

where $\varepsilon = \lambda \cdot \max_f |H(f)|^2$ and λ is the **regularization** parameter. With $\lambda = 0$ this reduces to exact division; small positive values stabilise the result near spectral nulls.

Listing 32: Regularised spectral deconvolution

```
use spectrograms::fft_deconvolve;

let recovered = fft_deconvolve(output_ne, kernel_ne, 1e-6)?;
```

12.3 Streaming Overlap-Save Convolution

For real-time or block-streaming applications (e.g., room-correction FIR filtering in an audio callback), repeatedly calling `fft_convolve` is wasteful because the impulse response h is fixed and its FFT could be cached. The `OverlapSaveConvolver` does exactly this.

Overlap-Save Method

At construction, the impulse response h of length K is zero-padded to the FFT size $N_{\text{FFT}} = \text{next_power_of_two}(B + K - 1)$ and transformed once. For each input block of length B :

1. Prepend the $N_{\text{FFT}} - B$ most recent samples (*history*) to the block to form a window of length N_{FFT} .
2. Compute an N_{FFT} -point C2c FFT.
3. Multiply element-wise by the cached H spectrum.
4. Inverse-transform.
5. Discard the first $K - 1$ alias-contaminated samples; the remaining B samples are the alias-free output.

Per-block cost: $O(B \log N_{\text{FFT}})$, compared to $O(BK)$ for direct convolution. **No heap allocation on the processing path.**

Listing 33: Streaming overlap-save with fixed block size

```
use spectrograms::OverlapSaveConvolver;
use std::num::NonZeroUsize;

// Build once from the impulse response.
let conv = OverlapSaveConvolver::new(&ir, NonZeroUsize::new(
    256).unwrap())?;

// In the audio callback (no allocation):
let mut output = vec![0.0f64; 256];
conv.process_block(&input_block, &mut output)?;
```

Key properties:

- Fixed block size set at construction; every call to `process_block` must provide exactly that many samples.
- Call `reset()` between independent streams to clear the inter-block history.
- Generic over `T: Sample` (works at both `f32` and `f64`).

12.4 Minimum-Phase Conversion

A *minimum-phase* filter has the same magnitude response as a given (typically linear-phase) FIR filter, but concentrates its energy at the start rather than the middle. This eliminates most of the filter’s latency—beneficial for live-processing and room-correction applications.

Real-Cepstrum Method

The library uses the homomorphic (real-cepstrum) algorithm:

1. Compute $H = \text{FFT}(h)$ on a heavily oversampled grid (default: $8 \times$ the IR length, rounded up to a power of two) to avoid cepstral time-aliasing.
2. Form the real log-magnitude: $\hat{H} = \log |H|$.
3. Apply an inverse FFT to obtain the real cepstrum c .
4. Window c : keep the DC and Nyquist taps unchanged, double all causal taps (indices 1 to $N/2 - 1$), and zero all anticausal taps (indices $N/2 + 1$ to $N - 1$). This is the discrete version of the Hilbert relation that derives minimum-phase angle from log-magnitude.
5. Compute $H_{\min} = \exp(\text{FFT}(c))$, then $h_{\min} = \text{Re}(\text{IFFT}(H_{\min}))$.
6. Truncate to the desired output length.

The magnitude response of $h_{\min}$ matches that of h to within the accuracy of the oversampled FFT; increasing the oversample factor reduces the residual magnitude error at the cost of a larger one-off transform.

Listing 34: Converting linear-phase FIR to minimum phase

```
use spectrograms::minimum_phase;

// A 64-tap linear-phase low-pass prototype.
let lin_phase_fir: Vec<f32> = /* ... */;
let mp_fir = minimum_phase(&lin_phase_fir)?;

// mp_fir has the same length and (nearly) the same
// magnitude response,
// but most energy is in the first few taps.
assert!(mp_fir[0].abs() >= mp_fir[mp_fir.len() - 1].abs());
```

When to use minimum phase. A linear-phase FIR of length L introduces a group delay of $(L - 1)/2$ samples. For a 1000-tap room correction filter at 48 kHz this is over 10 ms of latency. The minimum-phase equivalent reduces

this to near zero while preserving the equalisation curve. Use `minimum_phase_` with to control the output length (useful when only the leading portion of the minimum-phase response is needed) and the oversampling factor (higher values give a more accurate magnitude match).




Figure 20: Linear-phase vs. minimum-phase conversion of a 64-tap low-pass FIR filter ($f_c = 2$ kHz, $f_s = 22050$ Hz, Hamming window) via the real-cepstrum method. **Top-left:** impulse responses—the linear-phase filter (blue) is symmetric around its centre tap; the minimum-phase equivalent (orange) concentrates energy at the start. **Top-right:** magnitude responses—both filters share *identical* frequency responses; only the phase differs. **Bottom-left:** phase responses. **Bottom-right:** group delay—the linear-phase filter introduces $(L - 1)/2 \approx 1.45$ ms of constant latency; the minimum-phase design reduces this to near zero, enabling low-latency room correction and live processing.

13 Computational Optimizations

13.1 Plan-Based Computation

FFT planning involves:

1. Analyzing the transform size
2. Selecting optimal algorithm (radix-2, split-radix, Bluestein, etc.)
3. Precomputing twiddle factors: $W_N^k = e^{-j2\pi k/N}$
4. Allocating scratch buffers

For a single FFT, planning overhead is negligible. But for spectrograms with hundreds of frames, planning each FFT is wasteful.

Solution: Reusable plans

Listing 35: Reusable STFT plan with cached FFT and window

```
pub struct StftPlan {
    n_fft: usize,
    hop: usize,
    window: Vec<f64>,           // Pre-computed window
    fft: Box<dyn R2cPlan>,      // Reusable FFT plan
    fft_out: Vec<Complex<f64>>, // Reusable buffer
    frame: Vec<f64>,           // Reusable frame buffer
}

impl StftPlan {
    pub fn new(params: &SpectrogramParams) -> Result<Self> {
        let n_fft = params.stft().n_fft();
        let window = make_window(params.stft().window(),
            n_fft)?;

        // Create FFT plan once
        let mut planner = RealFftPlanner::new();
        let fft = planner.plan_r2c(n_fft)?;

        Ok(Self {
            n_fft,
            window,
            fft,
            fft_out: vec![Complex::new(0.0, 0.0); n_fft/2 +
1],
            frame: vec![0.0; n_fft],
        })
    }
}
```

```

    }

    // Reuse plan for each frame
    fn compute_frame(&mut self, samples: &[f64],
                    frame_idx: usize) -> Result<()> {
        // Window frame into self.frame buffer
        // ...

        // Reuse FFT plan and buffers
        self.fft.process(&self.frame, &mut self.fft_out)?;
        Ok(())
    }
}

```

Computational benefit: For 1000-frame spectrograms, plan reuse amortizes the planning overhead across all frames.

13.2 Sparse Matrix Operations

For filterbank multiplication $\mathbf{y} = H\mathbf{x}$ where H is sparse:

Dense: $\mathcal{O}(mn)$ for $m \times n$ matrix

Sparse: $\mathcal{O}(nnz)$ where nnz is number of non-zeros

For mel filterbanks with 80 bands and 257 FFT bins:

- Dense: $80 \times 257 = 20,560$ ops
- Sparse (~ 20 non-zeros/row): $80 \times 20 = 1,600$ ops

Row-wise sparse storage:

Listing 36: Sparse matrix layout and filterbank multiply

```

struct SparseMatrix {
    nrows: usize,
    ncols: usize,
    values: Vec<Vec<f64>>, // Non-zero values per row
    indices: Vec<Vec<usize>>, // Column indices per row
}

fn multiply_vec(&self, input: &[f64], out: &mut [f64]) {
    for (row_idx, (row_values, row_indices)) in
        self.values.iter().zip(&self.indices).enumerate() {
        let mut acc = 0.0;
        for (&value, &col_idx) in

```

```

        row_values.iter().zip(row_indices) {
            acc += value * input[col_idx];
        }
        out[row_idx] = acc;
    }
}

```

Cache efficiency: Row-wise format ensures sequential access to both input array and sparse values, maximizing cache hits.

13.3 Design Patterns

Minimize allocations in hot loops:

Bad: Allocate per iteration

Listing 37: Per-frame allocation — wasteful hot-loop pattern

```

for frame in 0..n_frames {
    let mut windowed = vec![0.0; n_fft]; // Allocation!
    // ... process frame
}

```

Good: Reuse workspace

Listing 38: Pre-allocated workspace reused across frames

```

let mut workspace = vec![0.0; n_fft]; // Single allocation
for frame in 0..n_frames {
    // Reuse workspace
    fill_frame(&mut workspace, samples, frame);
    // ... process frame
}

```

Workspace pattern:

Listing 39: Full workspace struct with resizable buffers

```

struct Workspace {
    frame: Vec<f64>, // Windowed frame buffer
    fft_out: Vec<Complex<f64>>, // FFT output buffer
    spectrum: Vec<f64>, // Power spectrum buffer
    mapped: Vec<f64>, // Filterbank output buffer
}

impl Workspace {
    fn ensure_sizes(&mut self, n_fft: usize,
                    out_len: usize, n_bins: usize) {

```

```

        self.frame.resize(n_fft, 0.0);
        self.fft_out.resize(out_len, Complex::new(0.0, 0.0))
    ;

    self.spectrum.resize(out_len, 0.0);
    self.mapped.resize(n_bins, 0.0);
}
}

```

13.4 Memory Layout and Cache Efficiency

Row-major vs. column-major:

The library uses row-major layout (C-style) for compatibility with most FFT libraries:

Listing 40: Asserting row-major layout for FFT compatibility

```

// Row-major: rows are contiguous
let data = Array2::::zeros((nrows, ncols));
assert!(data.is_standard_layout()); // row-major

```

For column-wise operations (processing spectrogram frames), this can cause cache misses. But the benefit of contiguous memory for FFT outweighs this.

Alignment: Modern FFT libraries (FFTW, RealFFT) benefit from aligned memory. The library ensures contiguous slices are passed to FFT routines.

13.5 FFTW Integration

FFTW (*Fastest Fourier Transform in the West*) [10] is considered the gold standard for FFT performance. The library supports both:

- **RealFFT:** Pure Rust, portable, good performance
- **FFTW:** C library, excellent performance, requires system installation

FFTW wisdom:

FFTW learns optimal algorithms through runtime measurements. Plans can be saved and reused:

Listing 41: FFTW planning modes from ESTIMATE to PATIENT

```

// FFTW planning flags:
// ESTIMATE: Fast planning, ok performance
// MEASURE: Slower planning, better performance

```



```
// PATIENT: Very slow planning, best performance
let plan = fftw::plan_r2c_1d(n_fft, FFTW_MEASURE);
```

For repeated transforms of the same size, MEASURE or PATIENT modes can provide better performance than ESTIMATE.

14 Numerical Considerations

14.1 Floating-Point Precision

The library defaults to **f64** (double precision), and double precision is the recommended choice for offline analysis:

- 53 bits of precision (≈ 16 decimal digits)
- Dynamic range: $\approx 10^{-308}$ to 10^{308}
- Ample headroom for all audio processing needs

As described in Section 1.3, the core algorithms are generic over the scalar type and may also be instantiated at **f32**.

When to prefer f64. Double precision is advantageous when audio processing requires:

- Accurate accumulation over many samples (STFT frames)
- Stable recursive filters (if implemented)
- Wide dynamic range for decibel conversions

When to prefer f32. Single precision halves memory traffic for memory-bound workloads, matches machine-learning pipelines that train in single precision, and suits embedded or **no_std** targets with single-precision floating-point units. **f64** remains the default because it ensures numerical stability without performance concerns on modern CPUs.

14.2 Normalization Conventions

Different conventions exist for FFT normalization. The library uses:

Forward FFT: No normalization

$$X[k] = \sum_{n=0}^{N-1} x[n] e^{-j2\pi kn/N} \quad (96)$$

Inverse FFT: Divide by N

$$x[n] = \frac{1}{N} \sum_{k=0}^{N-1} X[k] e^{j2\pi kn/N} \quad (97)$$

This ensures $\text{IFFT}(\text{FFT}(x)) = x$ exactly.

Alternative conventions (used by some libraries):

- Both directions scaled by $1/\sqrt{N}$ (symmetric)
- Inverse scaled by $1/N$, forward by N (rare)

14.3 Dynamic Range and Epsilon Thresholds

Logarithm floor:

To avoid $\log(0) = -\infty$:

Listing 42: Log floor with epsilon to avoid negative infinity

```
let epsilon = 1e-10;
let db = 10.0 * (power.max(epsilon)).log10();
```

Choosing ϵ :

- Too large: Loss of quiet signals
- Too small: Numerical instability
- Sweet spot: 10^{-10} to 10^{-8} for f64

Sparse matrix threshold:

For considering values “zero”:

Listing 43: Epsilon threshold for sparse matrix storage

```
if value.abs() > 1e-10 {
    store(value);
}
```

This is much larger than machine epsilon ($\approx 2.2 \times 10^{-16}$ for f64) to account for accumulated rounding errors.

15 Conclusion and Further Reading

15.1 Summary of Key Concepts

This manual covered:

Core algorithms:

- DFT and FFT (Cooley-Tukey algorithm)
- Real-to-complex FFT with Hermitian symmetry
- 2D FFT via row-column decomposition
- STFT and time-frequency uncertainty

Signal processing methods:

- Window functions (Hanning, Hamming, Blackman, Kaiser, Gaussian)
- Amplitude scaling (power, magnitude, decibels)
- Convolution theorem for efficient filtering

Perceptual scales:

- Mel scale for speech processing
- ERB scale (Glasberg & Moore) for auditory modeling
- Logarithmic frequency for music
- CQT with variable-length kernels

Advanced features:

- MFCCs via DCT for speech recognition
- Chroma features for music analysis
- Gammatone filters for perceptual accuracy
- Binaural spectrograms: ITD, IPD, ILD, ILR for spatial audio analysis
- MDCT/IMDCT with sine-window perfect reconstruction
- FFT convolution, deconvolution, and overlap-save streaming
- Minimum-phase conversion via the real-cepstrum method

Optimizations:

- Plan-based computation

- Sparse matrix operations
- Workspace patterns (zero allocation in hot loops)
- Cache-efficient memory layouts
- Optional Rayon parallelisation (**rayon** feature)

15.2 Applications and Extensions

The techniques in this manual enable:

- Speech recognition and speaker identification
- Music information retrieval
- Audio classification and tagging
- Sound event detection
- Acoustic scene analysis
- Image denoising and enhancement
- Texture analysis
- Pattern recognition

Software:

- FFTW: www.fftw.org
- Librosa (Python): librosa.org
- Numpy (Python): numpy.org
- SciPy (Python): scipy.org
- This library: github.com/jmg049/Spectrograms

16 Python API Reference

16.1 Introduction

This section provides a complete reference for the Python API, demonstrating how to apply the mathematical concepts from previous sections in practice. All code examples use `samples` (audio data as `ndarray[float64]`) and `image` (image data as `ndarray[float64]`).

Installation and Imports

Listing 44: Installing the spectrograms package

```
pip install spectrograms
```

Listing 45: Importing numpy and spectrograms

```
import numpy as np
import spectrograms as sg
```

API Organization

The API consists of four components:

- **Parameter Classes:** Configuration objects
- **Computation Functions:** Single-use convenience functions
- **Planner Classes:** Reusable plans for batch processing
- **Result Objects:** Rich results with metadata

All computation functions release Python's GIL for parallel processing.

Error Handling

Listing 46: Catching spectrogram-specific exceptions

```
try:
    spec = sg.compute_mel_power_spectrogram(samples, params,
                                             mel_params)
except sg.InvalidInputError as e:
    print(f"Invalid parameters: {e}")
except sg.DimensionMismatchError as e:
    print(f"Array size mismatch: {e}")
except sg.FFTBackendError as e:
    print(f"FFT computation error: {e}")
except sg.SpectrogramError as e:
    print(f"Library error: {e}")
```

Basic FFT Operations

1D FFT

Compute the Discrete Fourier Transform (Equation (1)):

Listing 47: Computing a 1D FFT with Hermitian output

```
n_fft: int = 512

# Compute FFT
spectrum: npt.NDArray[np.complex128] = sg.compute_fft(
    samples[:n_fft], n_fft)
print(f"Output shape: {spectrum.shape}") # (257,) due to
    Hermitian symmetry
print(f"Output dtype: {spectrum.dtype}") # complex128
```

Inverse FFT

Perfect reconstruction (Equation (3)):

Listing 48: Inverse FFT and reconstruction error check

```
spectrum: npt.NDArray[np.complex128] = sg.compute_fft(
    samples[:512], 512)
reconstructed: npt.NDArray[np.float64] = sg.compute_irfft(
    spectrum, 512)

error: float = np.max(np.abs(samples[:512] - reconstructed))
```

```
print(f"Reconstruction error: {error:.2e}") # < 1e-14
```

Power and Magnitude Spectra

From Section 5.1:

Listing 49: Power and magnitude spectra with windowing

```
# Power spectrum:  $|X[k]|^2$ 
power: npt.NDArray[np.float64] = sg.compute_power_spectrum(
    samples[:512], 512)

# Magnitude spectrum:  $|X[k]|$ 
magnitude: npt.NDArray[np.float64] = sg.
    compute_magnitude_spectrum(samples[:512], 512)

# With windowing (sec:windows)
power_windowed: npt.NDArray[np.float64] = sg.
    compute_power_spectrum(samples[:512], 512, window=
        WindowType.hanning)
```

Window Functions

Predefined Windows

Window functions from Section 3:

Listing 50: Selecting predefined window function types

```
# WindowType class methods
rect: sg.WindowType = sg.WindowType.rectangular
hann: sg.WindowType = sg.WindowType.hanning
hamming: sg.WindowType = sg.WindowType.hamming
blackman: sg.WindowType = sg.WindowType.blackman()

# Use with compute functions
power: npt.NDArray[np.float64] = sg.compute_power_spectrum(
    samples[:512], 512, window=hann)
```

Parametric Windows

Kaiser (Section 3.7) and Gaussian (Section 3.8):

Listing 51: Parametric Kaiser and Gaussian windows

```
# Kaiser window with beta parameter
kaiser: sg.WindowType = sg.WindowType.kaiser(beta=8.6)
power_kaiser: npt.NDArray[np.float64] = sg.
    compute_power_spectrum(samples[:512], 512, window=kaiser
    )

# Gaussian window with std
gaussian: sg.WindowType = sg.WindowType.gaussian(std=0.4)
power_gaussian: npt.NDArray[np.float64] = sg.
    compute_power_spectrum(samples[:512], 512, window=
    gaussian)
```

STFT and Spectrograms

Configuration

STFT parameters from Section 4.2:

Listing 52: Configuring STFT parameters and presets

```
# Manual configuration
window: sg.WindowType = sg.WindowType.hanning
stft: sg.StftParams = sg.StftParams(
    n_fft=512,
    hop_size=160,
    window=window,
    centre=True
)
params: sg.SpectrogramParams = sg.SpectrogramParams(stft,
    sample_rate=16000)

# Presets
params_speech: sg.SpectrogramParams = sg.SpectrogramParams.
    speech_default(16000)
params_music: sg.SpectrogramParams = sg.SpectrogramParams.
    music_default(44100)
```

Linear Spectrograms

Three amplitude representations (Section 5):

Listing 53: Power, magnitude, and dB linear spectrograms

```
params: sg.SpectrogramParams = sg.SpectrogramParams.
    speech_default(16000)

# Power:  $|X[k]|^2$ 
spec_power: sg.Spectrogram = sg.
    compute_linear_power_spectrogram(samples, params)

# Magnitude:  $|X[k]|$ 
spec_magnitude: sg.Spectrogram = sg.
    compute_linear_magnitude_spectrogram(samples, params)

# Decibels:  $10 \cdot \log_{10}(|X[k]|^2)$ 
db_params: sg.LogParams = sg.LogParams(floor_db=-80.0)
spec_db: sg.Spectrogram = sg.compute_linear_db_spectrogram(
    samples, params, db_params)

print(f"Shape: {spec_power.shape}") # (257, n_frames)
```

Spectrogram Result Object

Listing 54: Accessing metadata from a Spectrogram object

```
spec: sg.Spectrogram = sg.compute_linear_db_spectrogram(
    samples, params, db_params)

# Access as ndarray via buffer protocol
data: npt.NDArray[np.float64] = np.asarray(spec) # Shape: (
    n_bins, n_frames)

# Direct property access
frequencies: npt.NDArray[np.float64] = spec.frequencies
times: npt.NDArray[np.float64] = spec.times
n_bins: int = spec.n_bins
n_frames: int = spec.n_frames

# Methods
duration: float = spec.duration()
freq_range: tuple[float, float] = spec.frequency_range()
db_range: tuple[float, float] | None = spec.db_range()
```

Perceptual Frequency Scales

Mel Scale

Mel-scale spectrograms (Section 6.2):

Listing 55: Computing mel-scale spectrograms

```
mel_params: sg.MelParams = sg.MelParams(  
    n_mels=80,  
    f_min=0.0,  
    f_max=8000.0  
)  
  
# Three amplitude types  
mel_power: sg.Spectrogram = sg.compute_mel_power_spectrogram(  
    samples, params, mel_params)  
mel_magnitude: sg.Spectrogram = sg.  
    compute_mel_magnitude_spectrogram(samples, params,  
    mel_params)  
mel_db: sg.Spectrogram = sg.compute_mel_db_spectrogram(  
    samples, params, mel_params, db_params)  
  
print(f"Mel shape: {mel_power.shape}") # (80, n_frames)  
print(f"Mel frequencies: {mel_power.frequencies[:5]}")
```

ERB Scale

ERB-scale spectrograms (Section 6.3):

Listing 56: Computing ERB-scale spectrograms

```
erb_params: sg.ErbParams = sg.ErbParams(  
    n_filters=64,  
    f_min=50.0,  
    f_max=8000.0  
)  
  
erb_power: sg.Spectrogram = sg.compute_erb_power_spectrogram(  
    samples, params, erb_params)  
erb_magnitude: sg.Spectrogram = sg.  
    compute_erb_magnitude_spectrogram(samples, params,  
    erb_params)  
erb_db: sg.Spectrogram = sg.compute_erb_db_spectrogram(  
    samples, params, erb_params, db_params)
```

Logarithmic Hz Scale

Logarithmic frequency scale for music:

Listing 57: Computing log-Hz spectrograms for music

```
loghz_params: sg.LogHzParams = sg.LogHzParams(  
    n_bins=84,  
    f_min=55.0,  
    f_max=7040.0  
)  
  
loghz_power: sg.Spectrogram = sg.  
    compute_loghz_power_spectrogram(samples, params,  
    loghz_params)  
loghz_magnitude: sg.Spectrogram = sg.  
    compute_loghz_magnitude_spectrogram(samples, params,  
    loghz_params)  
loghz_db: sg.Spectrogram = sg.compute_loghz_db_spectrogram(  
    samples, params, loghz_params, db_params)
```

Constant-Q Transform

CQT (Section 7):

Listing 58: Computing CQT spectrograms

```
cqt_params: sg.CqtParams = sg.CqtParams(  
    bins_per_octave=12,  
    n_octaves=7,  
    f_min=55.0  
)  
  
cqt_power: sg.Spectrogram = sg.compute_cqt_power_spectrogram(  
    samples, params, cqt_params)  
cqt_magnitude: sg.Spectrogram = sg.  
    compute_cqt_magnitude_spectrogram(samples, params,  
    cqt_params)  
cqt_db: sg.Spectrogram = sg.compute_cqt_db_spectrogram(  
    samples, params, cqt_params, db_params)
```

Audio Features

Raw STFT

Complex-valued STFT matrix:

Listing 59: Raw complex STFT with magnitude and phase extraction

```
stft_matrix: npt.NDArray[np.complex128] = sg.compute_stft(
    samples, params.stft)
print(f"STFT shape: {stft_matrix.shape}") # (n_bins,
    n_frames)
print(f"STFT dtype: {stft_matrix.dtype}") # complex128

# Extract magnitude and phase
magnitude: npt.NDArray[np.float64] = np.abs(stft_matrix)
phase: npt.NDArray[np.float64] = np.angle(stft_matrix)
```

Inverse STFT

Reconstruct signal from STFT (Section 3.10):

Listing 60: Inverse STFT reconstruction and error verification

```
stft_matrix: npt.NDArray[np.complex128] = sg.compute_stft(
    samples, params.stft)

reconstructed: npt.NDArray[np.float64] = sg.compute_istft(
    stft_matrix,
    n_fft=params.stft.n_fft,
    hop_size=params.stft.hop_size,
    window=params.stft.window,
    centre=params.stft.centre
)

# Verify reconstruction
error: float = np.max(np.abs(samples[:len(reconstructed)] -
    reconstructed))
print(f"Error: {error:.2e}")
```

MFCC

Mel-Frequency Cepstral Coefficients (Section 8.1):

Listing 61: Computing MFCCs with liftering and DCT

```
mfcc_params: sg.MfccParams = sg.MfccParams(  
    n_mfcc=13,  
    use_dct_ortho=True,  
    lifter_coeff=22.0  
)  
  
mfccs: npt.NDArray[np.float64] = sg.compute_mfcc(  
    samples,  
    params.stft,  
    sample_rate=16000,  
    n_mels=40,  
    mfcc_params=mfcc_params  
)  
print(f"MFCC shape: {mfccs.shape}") # (13, n_frames)
```

Chromagram

Pitch class profiles (Section 8.2):

Listing 62: Computing pitch-class chromagram

```
chroma_params: sg.ChromaParams = sg.ChromaParams(  
    tuning=440.0,  
    f_min=32.7,  
    f_max=4186.0,  
    norm=ChromaNorm.l2  
)  
  
chroma: sg.Chromagram = sg.compute_chromagram(  
    samples,  
    params.stft,  
    sample_rate=16000,  
    chroma_params=chroma_params  
)  
  
# Access as ndarray (12 pitch classes, n_frames)  
chroma_data: npt.NDArray[np.float64] = np.asarray(chroma)  
  
pitch_classes = ['C', 'C#', 'D', 'D#', 'E', 'F', 'F#', 'G',  
    'G#', 'A', 'A#', 'B']
```

Performance Optimization

Plan Reuse

Reuse FFT plans for batch processing:

Listing 63: Reusing a mel-dB plan for batch processing

```
# Single computation
spec1: sg.Spectrogram = sg.compute_mel_db_spectrogram(
    samples, params, mel_params, db_params)

# Batch with plan reuse
planner: sg.SpectrogramPlanner = sg.SpectrogramPlanner()
plan = planner.mel_db_plan(params, mel_params, db_params)

signals = [samples1, samples2, samples3]
spectrograms = [plan.compute(s) for s in signals]
```

Available Plans

All spectrogram types have corresponding plans:

Listing 64: All spectrogram plan types from SpectrogramPlanner

```
planner = sg.SpectrogramPlanner()

# Linear (3 types)
linear_power_plan = planner.linear_power_plan(params)
linear_magnitude_plan = planner.linear_magnitude_plan(params)
linear_db_plan = planner.linear_db_plan(params, db_params)

# Mel (3 types)
mel_power_plan = planner.mel_power_plan(params, mel_params)
mel_magnitude_plan = planner.mel_magnitude_plan(params,
    mel_params)
mel_db_plan = planner.mel_db_plan(params, mel_params,
    db_params)

# ERB (3 types)
erb_power_plan = planner.erb_power_plan(params, erb_params)
erb_magnitude_plan = planner.erb_magnitude_plan(params,
    erb_params)
erb_db_plan = planner.erb_db_plan(params, erb_params,
    db_params)
```

```

# LogHz (3 types)
loghz_power_plan = planner.loghz_power_plan(params,
    loghz_params)
loghz_magnitude_plan = planner.loghz_magnitude_plan(params,
    loghz_params)
loghz_db_plan = planner.loghz_db_plan(params, loghz_params,
    db_params)

# CQT (3 types)
cqt_power_plan = planner.cqt_power_plan(params, cqt_params)
cqt_magnitude_plan = planner.cqt_magnitude_plan(params,
    cqt_params)
cqt_db_plan = planner.cqt_db_plan(params, cqt_params,
    db_params)

# All plans use same interface
spec = plan.compute(samples)

```

Frame-by-Frame Processing

Process individual frames without full spectrogram:

Listing 65: Frame-by-frame spectrogram computation

```

planner = sg.SpectrogramPlanner()
plan = planner.mel_power_plan(params, mel_params)

n_frames = (len(samples) + params.stft.n_fft - params.stft.
    hop_size) // params.stft.hop_size

for frame_idx in range(n_frames):
    frame = plan.compute_frame(samples, frame_idx) # 1D
    array (n_mels,)
    # Process frame...

```

Streaming Processing

Real-time frame extraction:

Listing 66: Streaming real-time frame extraction from buffer

```

planner = sg.SpectrogramPlanner()
plan = planner.mel_power_plan(params, mel_params)

buffer = np.array([], dtype=np.float64)

```



```

chunk_size = 160

for chunk in audio_stream:
    buffer = np.concatenate([buffer, chunk])

    while len(buffer) >= params.stft.n_fft:
        frame = plan.compute_frame(buffer, 0)
        # Process frame...
        buffer = buffer[params.stft.hop_size:]

```

2D FFT

Forward 2D FFT

2D FFT from Section 2.5:

Listing 67: Forward 2D FFT with half-spectrum output

```

image = np.load('test.png') # Shape: (height, width)

spectrum = sg.fft2d(image)
print(f"Spectrum shape: {spectrum.shape}") # (height, width
//2 + 1)
print(f"Spectrum dtype: {spectrum.dtype}") # complex64

```

Inverse 2D FFT

Listing 68: Inverse 2D FFT and reconstruction error

```

spectrum = sg.fft2d(image)
reconstructed = sg.ifft2d(spectrum, image.shape[1])

error = np.max(np.abs(image - reconstructed))
print(f"Reconstruction error: {error:.2e}")

```

2D Power and Magnitude Spectra

Listing 69: 2D power and magnitude spectra with log scale

```
power = sg.power_spectrum_2d(image)           # Shape: (height,
        width//2 + 1)
magnitude = sg.magnitude_spectrum_2d(image)

# Log scale for visualization
power_db = 10 * np.log10(power + 1e-10)
```

FFT Shift

Center DC component:

Listing 70: FFT shift to centre the DC component

```
power = sg.power_spectrum_2d(image)

# Shift DC to center
power_shifted = sg.fftshift(power)

# Reverse shift
power_restored = sg.ifftshift(power_shifted)
```

Batch 2D FFT

Listing 71: Batch 2D FFT via Fft2dPlanner plans

```
planner = sg.Fft2dPlanner()

# Forward plan
fft_plan = planner.fft2d_plan(nrows, ncols)
spectra = [fft_plan.forward(img) for img in images]

# Inverse plan
ifft_plan = planner.ifft2d_plan(nrows, ncols)
reconstructed = [ifft_plan.inverse(spec) for spec in spectra
                 ]

# Power spectrum plan
power_plan = planner.power_spectrum_plan(nrows, ncols)
powers = [power_plan.compute(img) for img in images]

# Magnitude spectrum plan
```

```
magnitude_plan = planner.magnitude_spectrum_plan(nrows,
ncols)
magnitudes = [magnitude_plan.compute(img) for img in images]
```

Image Filtering

Gaussian Blur

Listing 72: Gaussian blur via FFT convolution

```
kernel = sg.gaussian_kernel_2d(size=15, sigma=2.0)
blurred = sg.convolve_fft(image, kernel)
```

Frequency Filters

Listing 73: Low-pass, high-pass, and band-pass image filters

```
# Lowpass: smooth (keep low frequencies)
lowpass = sg.lowpass_filter(image, cutoff_fraction=0.1)

# Highpass: edges (keep high frequencies)
highpass = sg.highpass_filter(image, cutoff_fraction=0.1)

# Bandpass: specific frequency range
bandpass = sg.bandpass_filter(image, low_cutoff=0.1,
high_cutoff=0.5)
```

Edge Detection and Sharpening

Listing 74: FFT-based edge detection and image sharpening

```
# Edge detection
edges = sg.detect_edges_fft(image)

# Sharpening
sharpened = sg.sharpen_fft(image, amount=1.5)
```

Complete API Reference

Parameter Classes

WindowType:

Listing 75: WindowType constructors and variants

```
sg.WindowType.rectangular
sg.WindowType.hanning
sg.WindowType.hamming
sg.WindowType.blackman()
sg.WindowType.kaiser(beta: float)
sg.WindowType.gaussian(std: float)
```

StftParams:

Listing 76: StftParams constructor signature

```
stft = sg.StftParams(
    n_fft: int,
    hop_size: int,
    window: str | WindowType,
    center: bool = True
)
```

LogParams:

Listing 77: LogParams constructor for dB floor

```
db_params = sg.LogParams(floor_db: float)
```

SpectrogramParams:

Listing 78: SpectrogramParams constructors and presets

```
params = sg.SpectrogramParams(stft: StftParams, sample_rate:
    float)
params = sg.SpectrogramParams.speech_default(sample_rate:
    float)
params = sg.SpectrogramParams.music_default(sample_rate:
    float)
```

MelParams:

Listing 79: MelParams constructor with frequency range

```
mel_params = sg.MelParams(
    n_mels: int,
```

```

        f_min: float = 0.0,
        f_max: float | None = None
    )

```

ErbParams:

Listing 80: ErbParams constructor with gammatone option

```

erb_params = sg.ErbParams(
    n_bands: int,
    f_min: float = 50.0,
    f_max: float | None = None,
    use_gammatone: bool = True
)

```

LogHzParams:

Listing 81: LogHzParams constructor with octave resolution

```

loghz_params = sg.LogHzParams(
    n_bins: int,
    f_min: float,
    f_max: float,
    bins_per_octave: int = 12
)

```

CqtParams:

Listing 82: CqtParams constructor with sparsity threshold

```

cqt_params = sg.CqtParams(
    n_bins: int,
    f_min: float,
    bins_per_octave: int = 12,
    sparsity_threshold: float = 0.01
)

```

ChromaParams:

Listing 83: ChromaParams constructor and music preset

```

chroma_params = sg.ChromaParams(
    tuning_frequency: float = 440.0,
    n_chroma: int = 12,
    norm_type: str = "L2"
)
chroma_params = sg.ChromaParams.music_standard()

```

MfccParams:

Listing 84: MfccParams constructor and speech preset

```
mfcc_params = sg.MfccParams(
    n_mfcc: int = 13,
    use_energy: bool = True,
    lifter_coefficient: float = 22.0,
    normalize: bool = False
)
mfcc_params = sg.MfccParams.speech_standard()
```

Spectrogram Functions

All return Spectrogram objects with `.data`, `.frequencies`, `.times`, `.duration()`, `.frequency_range()`, `.db_range()`, however, they can also act as an `ndarray` via the array protocol with no additional function calls:

Listing 85: All spectrogram compute function signatures

```
# Linear
sg.compute_linear_power_spectrogram(samples, params)
sg.compute_linear_magnitude_spectrogram(samples, params)
sg.compute_linear_db_spectrogram(samples, params, db_params)

# Mel
sg.compute_mel_power_spectrogram(samples, params, mel_params)
sg.compute_mel_magnitude_spectrogram(samples, params, mel_params)
sg.compute_mel_db_spectrogram(samples, params, mel_params, db_params)

# ERB
sg.compute_erb_power_spectrogram(samples, params, erb_params)
sg.compute_erb_magnitude_spectrogram(samples, params, erb_params)
sg.compute_erb_db_spectrogram(samples, params, erb_params, db_params)

# LogHz
sg.compute_loghz_power_spectrogram(samples, params, loghz_params)
sg.compute_loghz_magnitude_spectrogram(samples, params, loghz_params)
sg.compute_loghz_db_spectrogram(samples, params, loghz_params, db_params)
```

```

# CQT
sg.compute_cqt_power_spectrogram(samples, params, cqt_params
)
sg.compute_cqt_magnitude_spectrogram(samples, params,
    cqt_params)
sg.compute_cqt_db_spectrogram(samples, params, cqt_params,
    db_params)

```

Audio Feature Functions

Listing 86: Audio feature compute function signatures

```

# STFT
stft_matrix = sg.compute_stft(samples, stft_params) #
    complex128, (n_bins, n_frames)

# Inverse STFT
samples = sg.compute_istft(stft_matrix, n_fft, hop_size,
    window, center)

# MFCC
mfccs = sg.compute_mfcc(samples, stft_params, sample_rate,
    n_mels, mfcc_params)

# Chromagram
chroma = sg.compute_chromagram(samples, stft_params,
    sample_rate, chroma_params)

```

FFT Functions

Listing 87: 1D FFT and spectrum compute function signatures

```

# 1D FFT
spectrum = sg.compute_fft(samples, n_fft) # complex128, (
    n_fft//2 + 1,)
samples = sg.compute_irfft(spectrum, n_fft)

# Power/magnitude
power = sg.compute_power_spectrum(samples, n_fft, window)
magnitude = sg.compute_magnitude_spectrum(samples, n_fft,
    window)

```

2D FFT Functions

Listing 88: 2D FFT function signatures and return types

```
# 2D FFT
spectrum = sg.fft2d(image) # complex64, (nrows, ncols//2 + 1)
image = sg.ifft2d(spectrum, output_ncols)

# Power/magnitude
power = sg.power_spectrum_2d(image)
magnitude = sg.magnitude_spectrum_2d(image)

# Shift
shifted = sg.fftshift(array)
restored = sg.ifftshift(shifted)
```

Image Processing Functions

Listing 89: Image processing function signatures

```
# Convolution
kernel = sg.gaussian_kernel_2d(size, sigma)
result = sg.convolve_fft(image, kernel)

# Filters
lowpass = sg.lowpass_filter(image, cutoff_fraction)
highpass = sg.highpass_filter(image, cutoff_fraction)
bandpass = sg.bandpass_filter(image, low_cutoff, high_cutoff)

# Edge detection and sharpening
edges = sg.detect_edges_fft(image)
sharpened = sg.sharpen_fft(image, amount)
```


References

- [1] Nasir Ahmed, T. Natarajan, and Kamisetty R Rao. “Discrete cosine transform”. In: *IEEE transactions on Computers* 100.1 (2006), pp. 90–93.
- [2] Jonathan Allen. “Short term spectral analysis, synthesis, and modification by discrete Fourier transform”. In: *IEEE transactions on acoustics, speech, and signal processing* 25.3 (2003), pp. 235–238.
- [3] Ralph Beebe Blackman and John W Tukey. “The measurement of power spectra from the point of view of communications engineering”. In: *Bell System Technical Journal* 37.1 (1958), pp. 185–282.
- [4] AP Blozinski. “On a convolution theorem for $L(p, q)$ spaces”. In: *Transactions of the American Mathematical Society* 164 (1972), pp. 255–265.
- [5] Judith C Brown. “Calculation of a constant Q spectral transform”. In: *The Journal of the Acoustical Society of America* 89.1 (1991), pp. 425–434.
- [6] Judith C Brown and Miller S Puckette. “An efficient algorithm for the calculation of a constant Q transform”. In: *The Journal of the Acoustical Society of America* 92.5 (1992), pp. 2698–2701.
- [7] James W Cooley and John W Tukey. “An algorithm for the machine calculation of complex Fourier series”. In: *Mathematics of computation* 19.90 (1965), pp. 297–301.
- [8] Steven Davis and Paul Mermelstein. “Comparison of parametric representations for monosyllabic word recognition in continuously spoken sentences”. In: *IEEE transactions on acoustics, speech, and signal processing* 28.4 (1980), pp. 357–366.
- [9] Harvey Fletcher. “Auditory patterns”. In: *Reviews of modern physics* 12.1 (1940), p. 47.
- [10] Matteo Frigo and Steven G Johnson. “The design and implementation of FFTW3”. In: *Proceedings of the IEEE* 93.2 (2005), pp. 216–231.
- [11] Takuya Fujishima. “Realtime chord recognition of musical sound: A system using common lisp music”. In: *Proceedings of the International Computer Music Conference*. 1999, pp. 464–467.
- [12] Dennis Gabor. “Theory of communication. Part 1: The analysis of information”. In: *Journal of the Institution of Electrical Engineers-Part III: Radio and Communication Engineering* 93.26 (1946), pp. 429–441.
- [13] Carl Friedrich Gauss. “Nachlass: Theoria interpolationis methodo nova tractata”. In: *Carl Friedrich Gauss Werke* 3 (1866). Original work from 1805, pp. 265–327.
- [14] Brian R Glasberg and Brian CJ Moore. “Derivation of auditory filter shapes from notched-noise data”. In: *Hearing research* 47.1-2 (1990), pp. 103–138.

- [15] Fredric J Harris. “On the use of windows for harmonic analysis with the discrete Fourier transform”. In: *Proceedings of the IEEE* 66.1 (1978), pp. 51–83.
- [16] James F Kaiser. “Nonrecursive digital filter design using window function”. In: *Proceedings of the 1974 IEEE International Symposium on Circuits and Systems*. IEEE. 1974, pp. 20–23.
- [17] Brian CJ Moore and Brian R Glasberg. “Suggested formulae for calculating auditory-filter bandwidths and excitation patterns”. In: *The journal of the Acoustical Society of America* 74.3 (1983), pp. 750–753.
- [18] Meinard Müller, Frank Kurth, and Michael Clausen. “Audio Matching via Chroma-Based Statistical Features.” In: *ISMIR*. Vol. 2005. 2005, p. 6.
- [19] Alan V Oppenheim. *Discrete-time signal processing*. Pearson Education India, 1999.
- [20] Lawrence R Rabiner and Biing-Hwang Juang. “Fundamentals of speech recognition”. In: Prentice Hall, 1993.
- [21] Stanley Smith Stevens, John Volkman, and Edwin B Newman. “A scale for the measurement of the psychological magnitude pitch”. In: *The journal of the Acoustical Society of America* 8.3 (1937), pp. 185–190.

This manual documents the mathematical foundations and computational implementation of the **spectrograms** library, a precision-generic Rust library for FFT-based audio and image analysis. It covers the Discrete Fourier Transform and fast FFT algorithms; windowed spectral analysis via the Short-Time Fourier Transform; perceptual frequency scales including the Mel scale, Equivalent Rectangular Bandwidth, and the Constant-Q Transform; advanced audio features such as Mel-frequency cepstral coefficients, chromagrams, and binaural analysis; two-dimensional Fourier methods for image filtering and edge detection; and numerical precision for single- and double-precision computation. Every algorithm is given in closed-form notation alongside its Rust and Python implementation.

WHAT'S INSIDE

- Fourier foundations — DFT, FFT, Hermitian symmetry, 2D transforms
- Window functions — Hann, Hamming, Blackman, Kaiser, Gaussian
- The STFT and spectrograms — power, magnitude, decibel scaling
- Perceptual scales — mel, ERB, log-frequency, constant-Q
- Audio features — MFCC, chroma, gammatone, binaural ITD/IPD/ILD
- Image processing via 2D FFT — blur, sharpening, edge detection
- MDCT, FFT convolution, and streaming overlap-save